

Designing Parallel Embedded System on Chip with Xilinx technology for Real Time Machine Learning and Analytics

Giulio Corradi

Chief of the Industrial safety/networking
and power systems control Division @ Xilinx Inc.



Giulio Corradi brings 25 year of experience of management, software engineering and development of ASICs and FPGA in industrial, automation and medical systems specifically in the field of control and communication for the major corporations. In 2006 Giulio joined Xilinx @ Munich as a chief engineer.

Friday 7^o November 2018, ??? AM
@Engineering Faculty, UNIMORE



UNIMORE

UNIVERSITÀ DEGLI STUDI DI
MODENA E REGGIO EMILIA



Field-Programmable Gate Arrays

Paolo Burgio

paolo.burgio@unimore.it



Outline

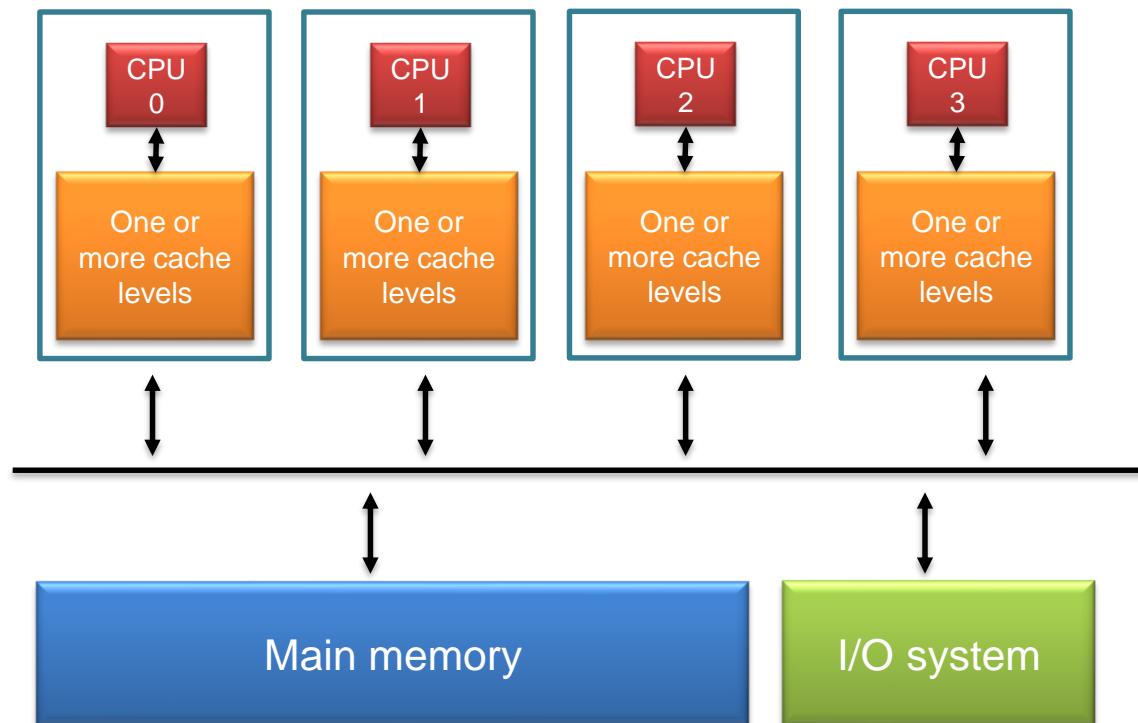
- ✓ Introduction to FPGAs
- ✓ How to use them
- ✓ Heterogeneous programming
- ✓ FPGA-based heterogeneous programming
- ✓ How to program it
- ✓ Xilinx: now and soon...



The world, till now

- ✓ (A) Symmetric multi-processing
 - Single or multi-core

Can be 1 bus, N
busses, or any
network






Reconfigurable Hardware

“Reconfigurable computing is intended to fill the gap between hardware and software, achieving potentially much higher performance than software, while maintaining a higher level of flexibility than hardware

(K. Compton and S. Hauck, *Reconfigurable Computing: a Survey of Systems and Software*, 2002)

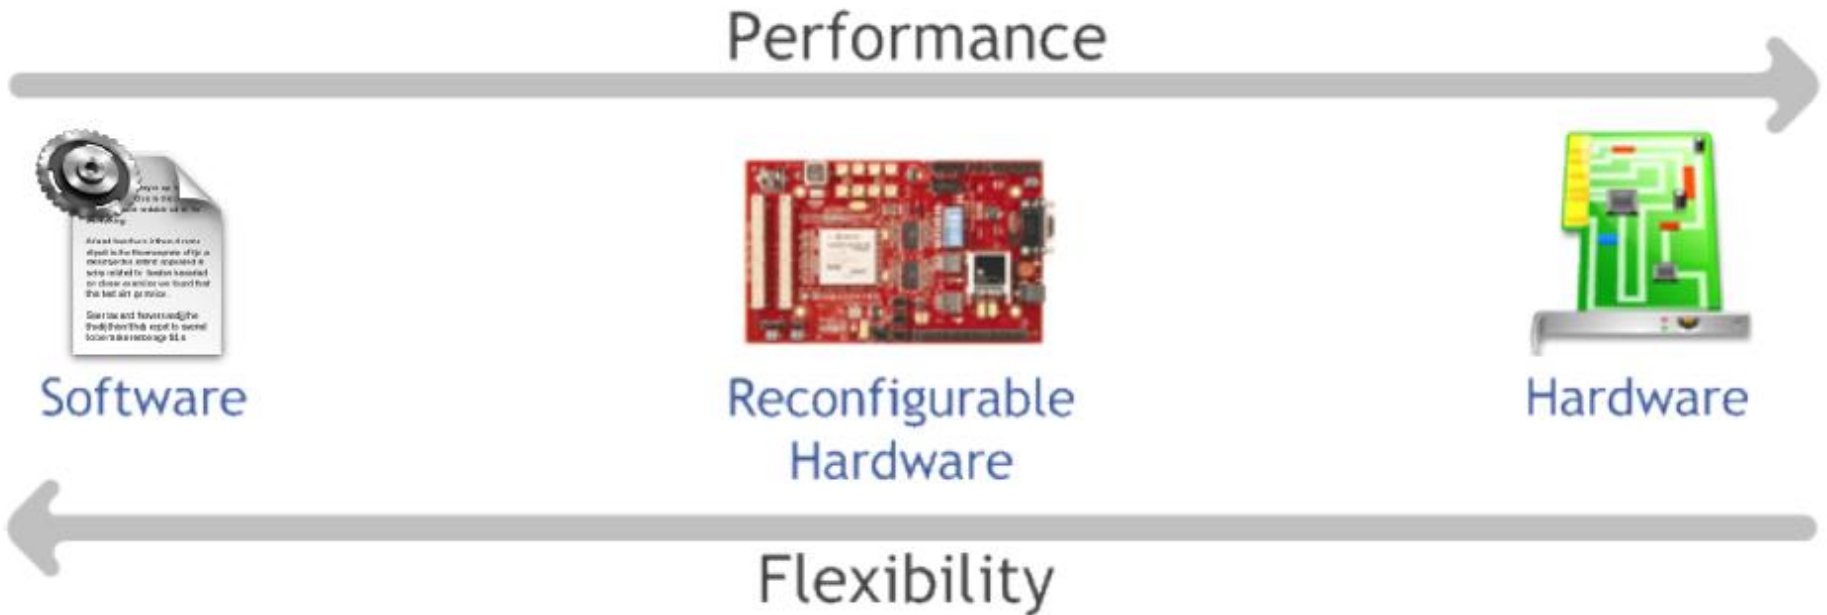




...eh?

We are used to have

- ✓ On one (left) side, full programmable artifacts (software)
 - Run on single or multi-cores, designed for General Purpose computing
- ✓ On one (right) side, hardware blocks to perform specific (subset of) operations
 - DSPs, co-processors





What if?

...we had a "sea" of hardware blocks that we can program as we want

- ✓ We can build cores
- ✓ We can build co-processors
- ✓ We can build what we want

What would you use them?

- ✓ For prototyping!


Hardware development process is long and cumbersome

- ✓ Imagine a full-fledged cores
- ✓ Typically, years of development
- ✓ You can "try, and see whether it works"



History of reconfigurable echnologies

- ✓ Logic gates (1950s-60s)
- ✓ Regular structures for two-level logic (1960s-70s)
 - Muxes and decoders, PLAs
- ✓ Programmable sum-of-products arrays (1970s-80s)
 - PLDs, complex PLDs
- ✓ Programmable gate arrays (1980s-90s)
 - densities high enough to permit entirely new class of application, e.g., prototyping, emulation, acceleration



trend toward
higher levels
of integration



Field-programmable gate arrays

- ✓ *"A field-programmable gate array (FPGA) is an integrated circuit designed to be **configured** by a customer or a designer **after manufacturing**."*
- ✓ Traditionally used for **prototyping**
 - Takes minutes vs years for "real" hardware
- ✓ Tech has evolved so they are actively used in production settings
 - Less powerful, yet more energy-efficient than a GPU
 - **Way** more flexible

Integrated into System-on-chips

- ✓ As reconfigurable accelerator
- ✓ We'll see later...




FPGAs

- ✓ Logic blocks
 - to implement combinational and sequential logic
- ✓ Interconnect
 - wires to connect inputs and outputs to logic blocks
- ✓ I/O blocks
 - special logic blocks at periphery of device for external connections


Key questions:

- ✓ how to make logic blocks programmable?
- ✓ how to connect the wires?
- ✓ after the chip has been fabbed





Commercial FPGA companies

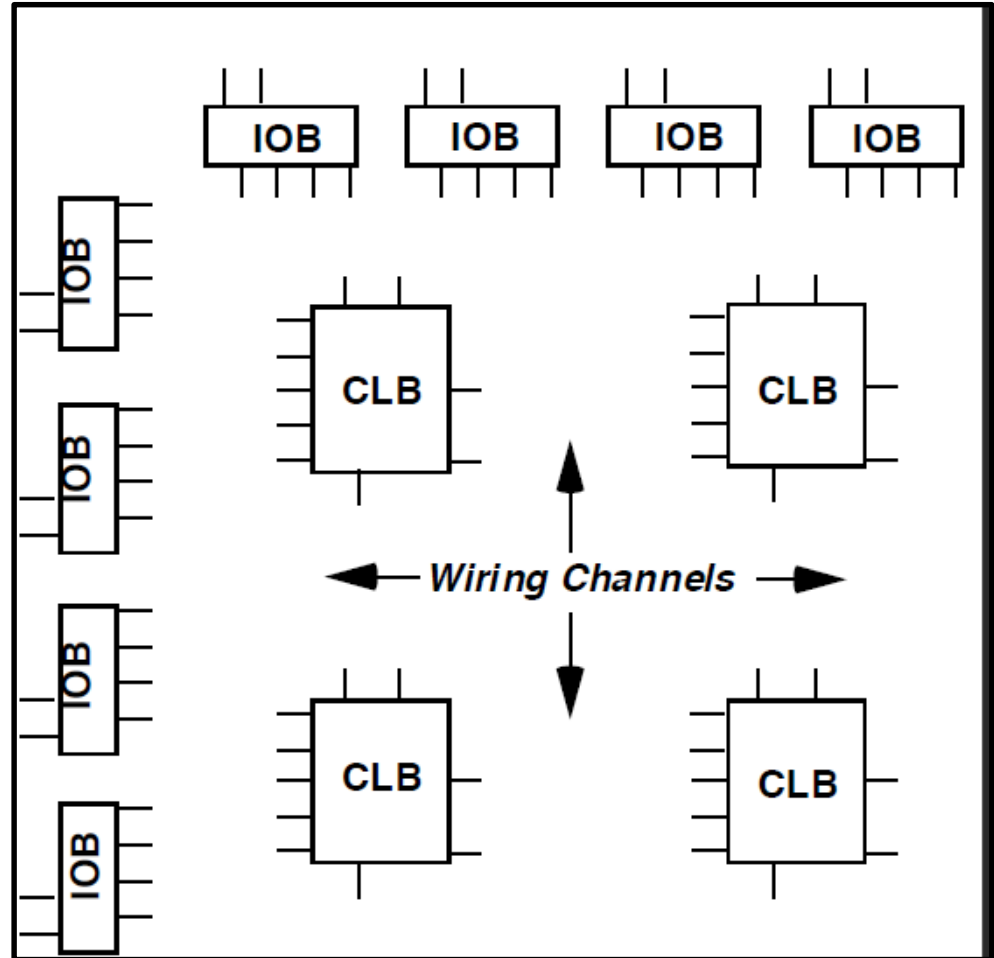




(Xilinx) Programmable Gate Arrays

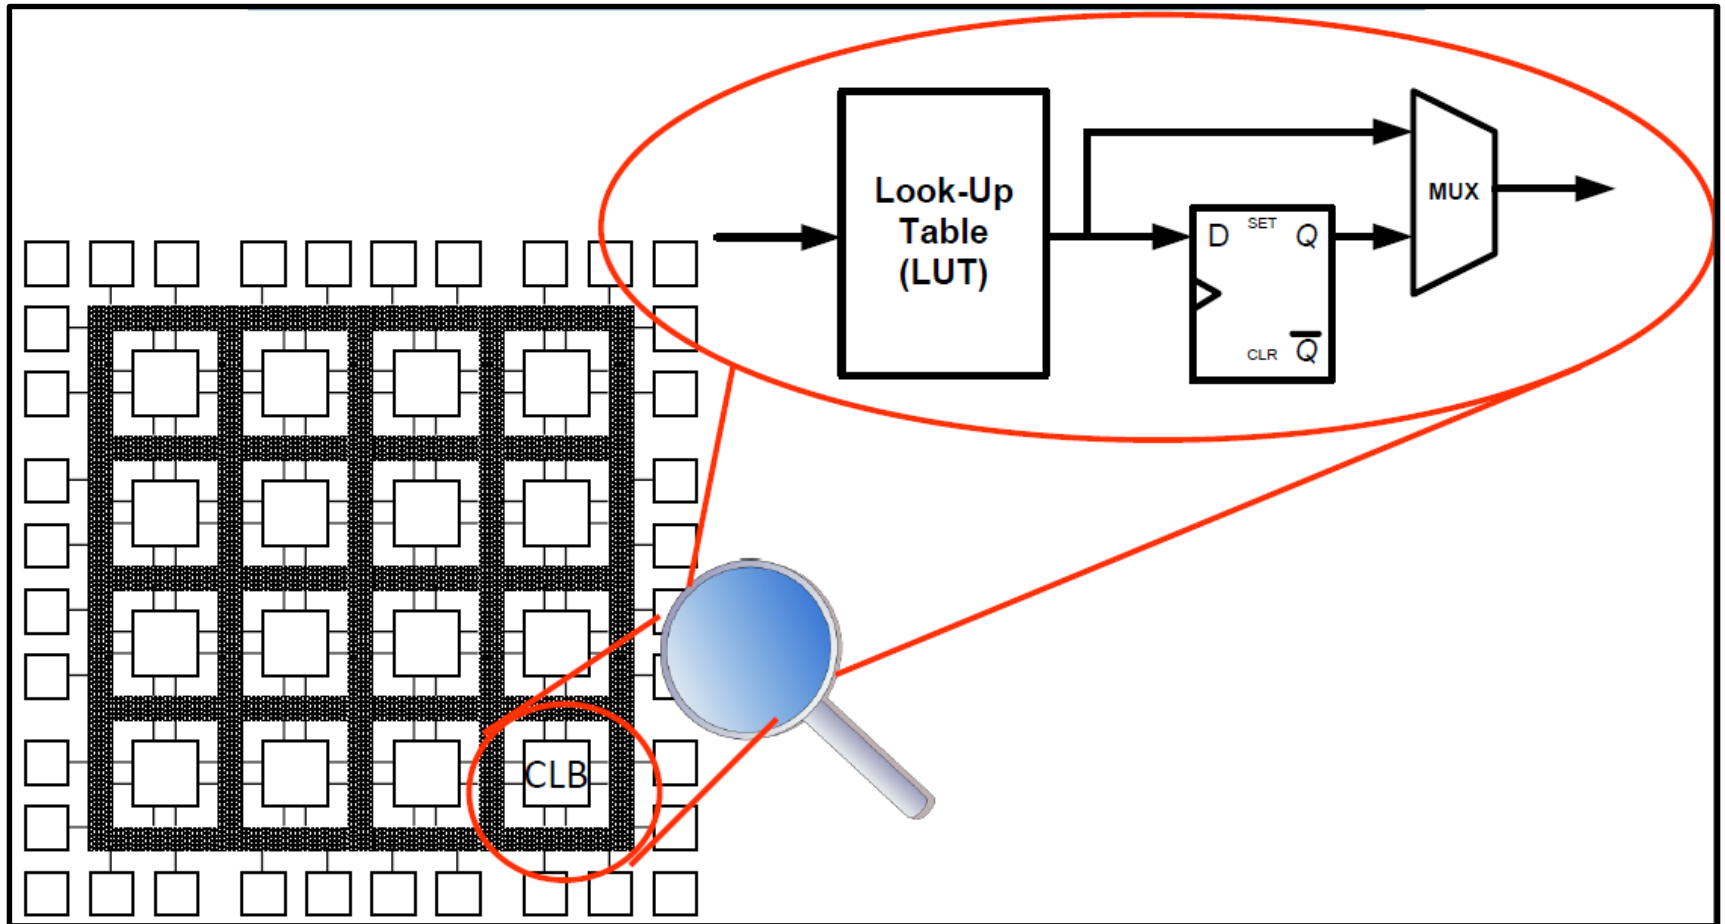
CLB - Configurable Logic Block

- ✓ Built-in fast carry logic
- ✓ Can be used as memory
- ✓ Three types of routing
 - direct
 - general-purpose
 - long lines of various lengths
- ✓ RAM-programmable
 - can be reconfigured





Simplified CLB Structure






LookUp Tables


LUT contains Memory Cells to implement small logic functions

- ✓ Each cell holds '0' or '1'
- ✓ Programmed with outputs of Truth Table
- ✓ Inputs select content of one of the cells as output

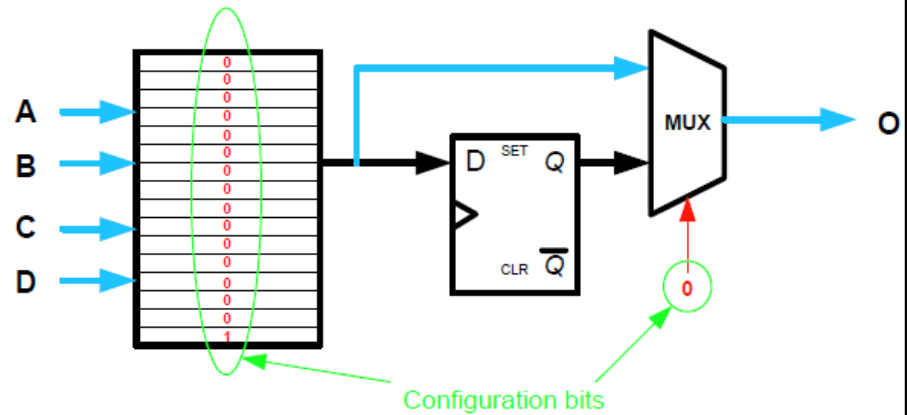




Example: 4-input AND gate

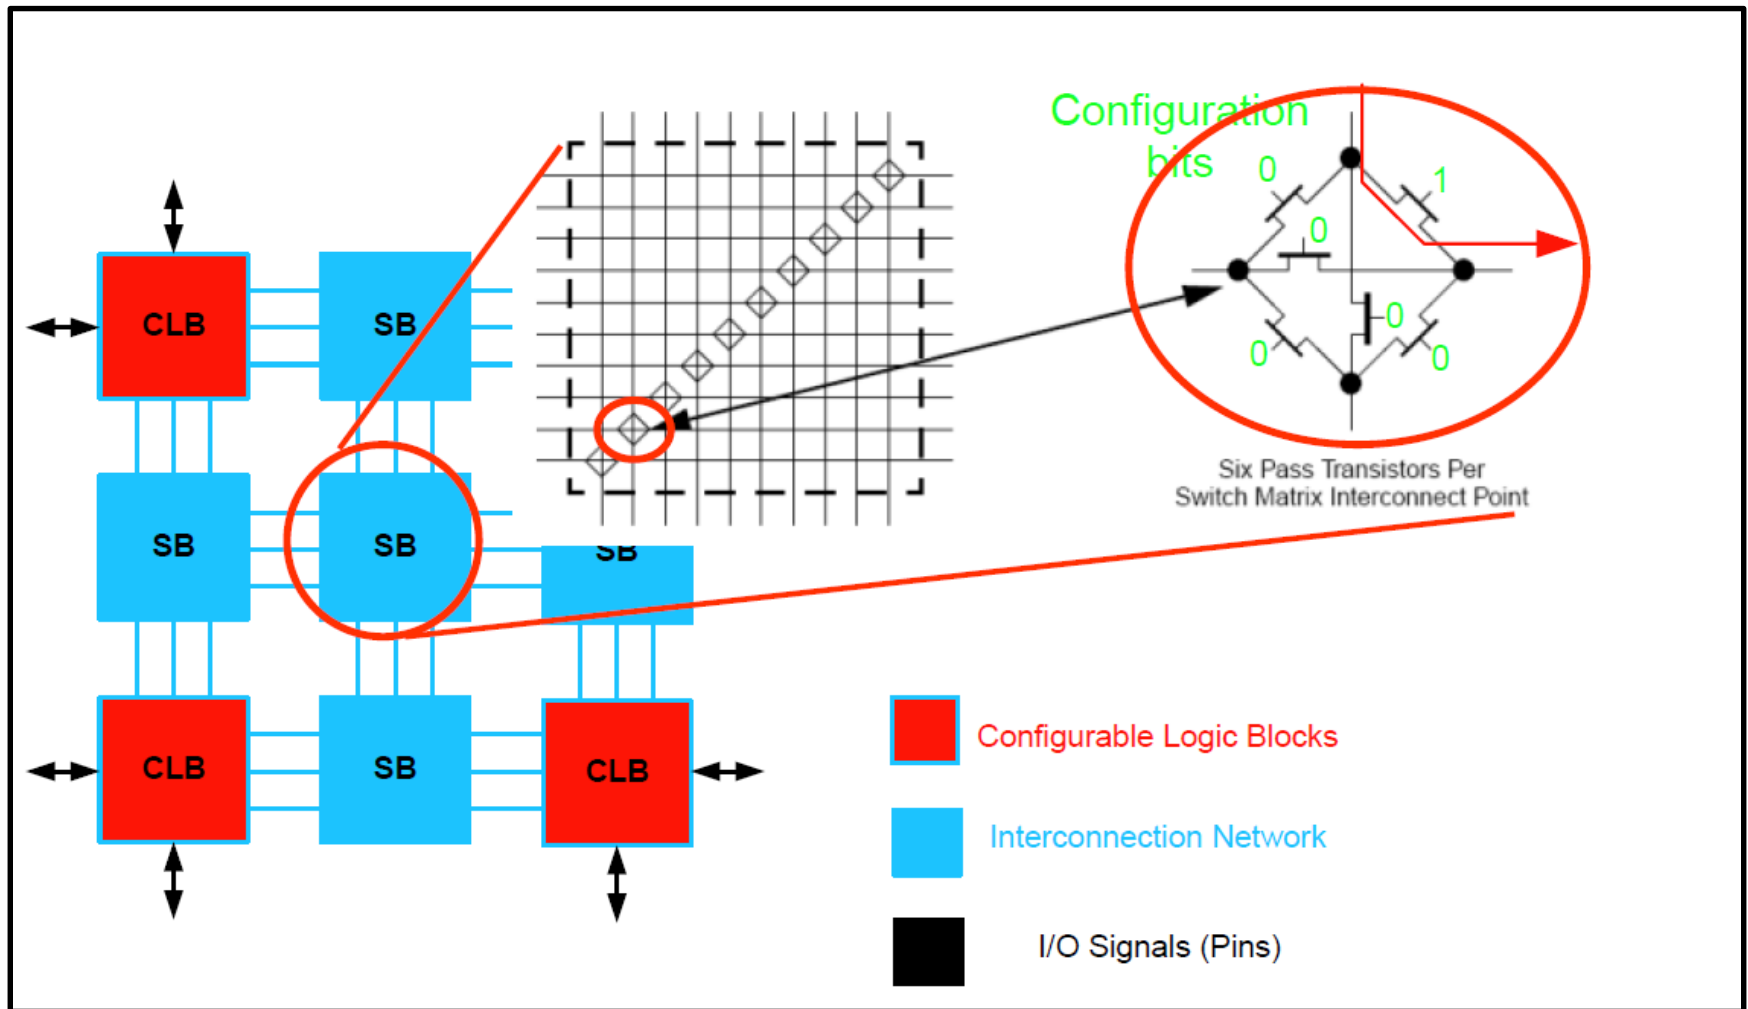


A	B	C	D	O
0	0	0	0	0
0	0	0	1	0
0	0	1	0	0
0	0	1	1	0
0	1	0	0	0
0	1	0	1	0
0	1	1	0	0
0	1	1	1	0
1	0	0	0	0
1	0	0	1	0
1	0	1	0	0
1	0	1	1	0
1	1	0	0	0
1	1	0	1	0
1	1	1	0	0
1	1	1	1	1





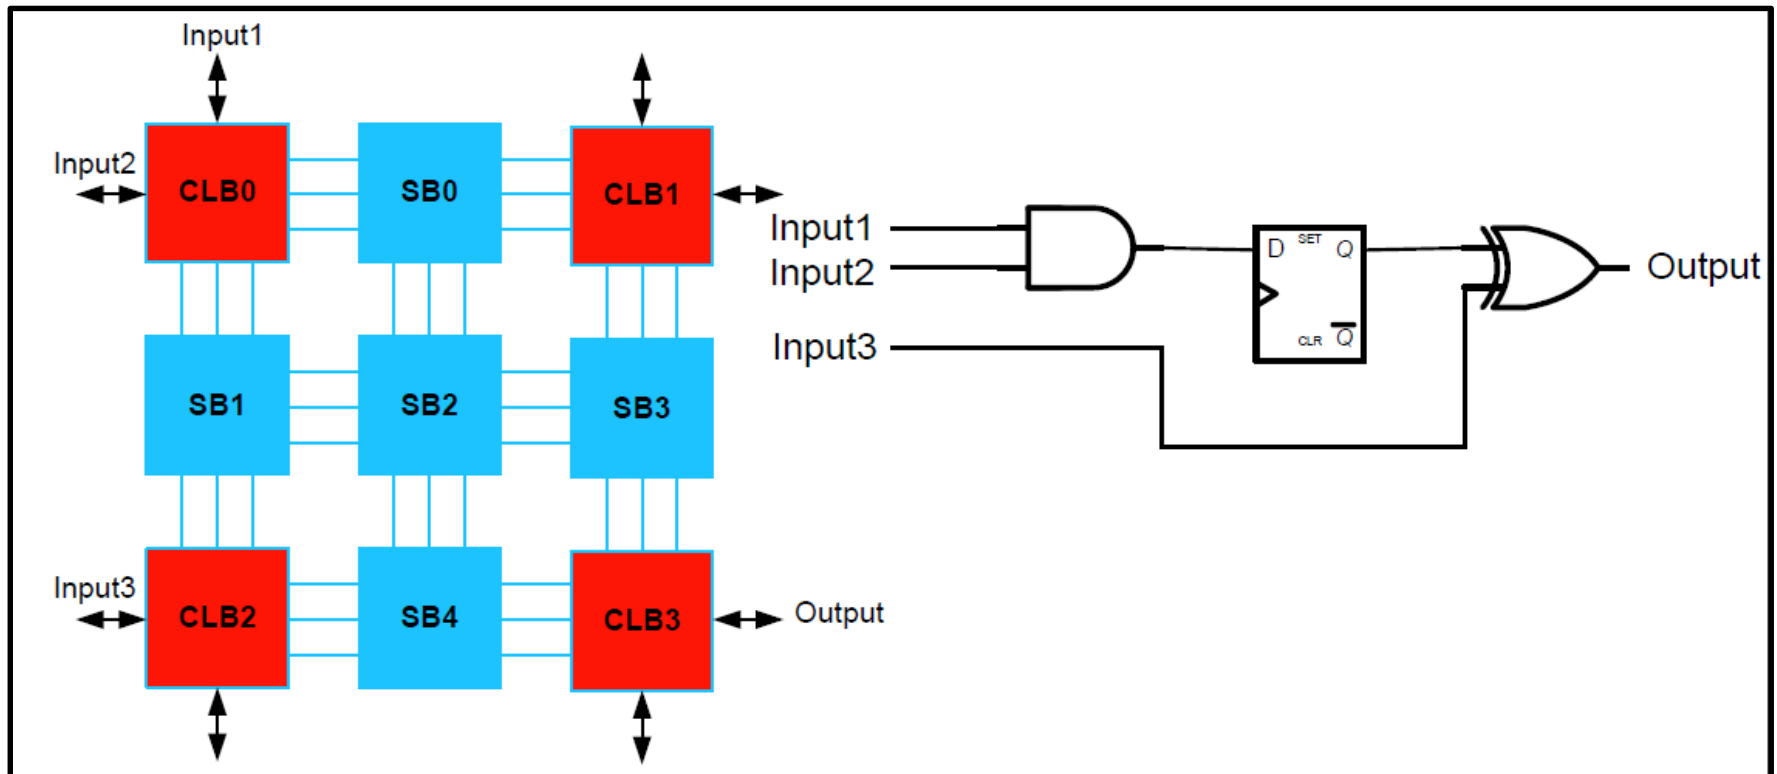
Interconnection Network






Example

- ✓ Determine the configuration bits for the following circuit implementation in a 2x2 FPGA, with I/O constraints as shown in the following figure. Assume 2-input LUTs in each CLB



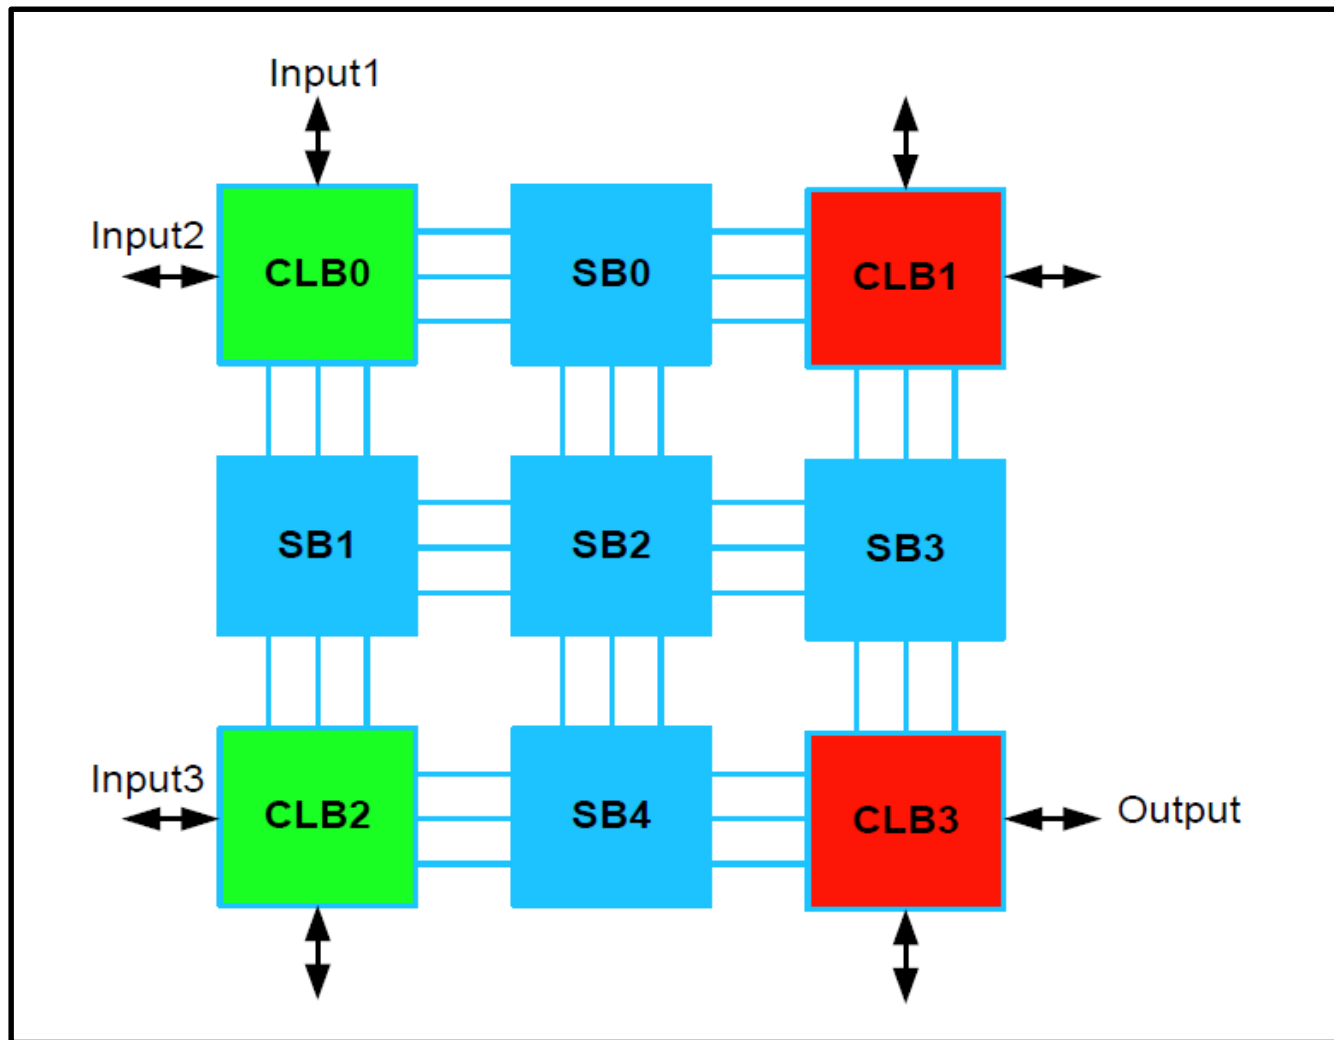


Configure CLB



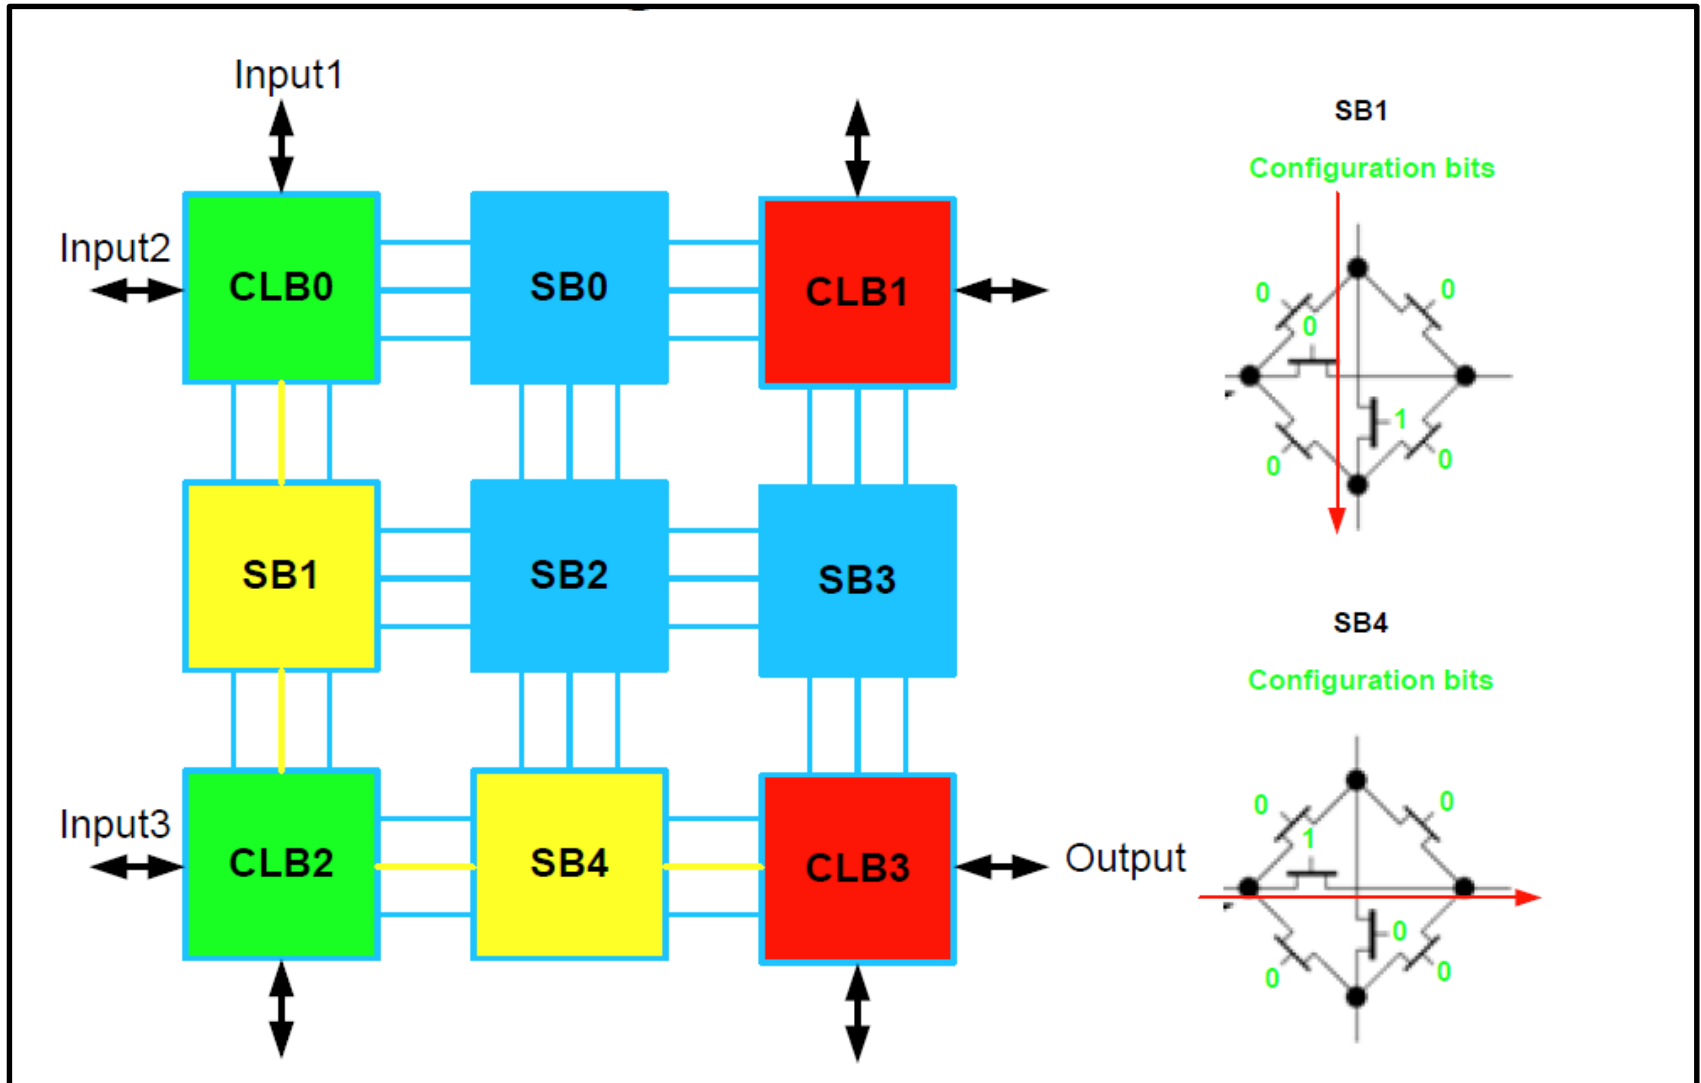


Placement: select CLBs





Routing: Select path





Configuration Bitstream

- ✓ The configuration bitstream must include ALL CLBs and SBs, even unused ones
- ✓ CLB0: 00011
- ✓ CLB1: ??????
- ✓ CLB2: 01100
- ✓ CLB3: XXXXX
- ✓ SB0: 000000
- ✓ SB1: 000010
- ✓ SB2: 000000
- ✓ SB3: 000000
- ✓ SB4: 000001



Some Definitions

- ✓ **Object Code (aka "Bitstream")**: the executable active physical (either HW or SW) implementation of a given functionality
- ✓ **Core**: a specific representation of a functionality. It is possible, for example, to have a core described in VHDL, in C or in an intermediate representation (e.g. a DFG)
- ✓ **IP-Core**: a core described using a HD Language combined with its communication infrastructure (i.e. the bus interface)
- ✓ **Reconfigurable Functional Unit**: an IP-Core that can be plugged and/or unplugged at runtime in an already working architecture
- ✓ **Reconfigurable Region**: a portion of the device area used to implement a reconfigurable core



Computer-Aided Design

Can't design FPGAs by hand!

- ✓ way too much logic to manage, hard to make changes
- ✓ **Hardware description languages** (HDL), es: Verilog, VHDL
 - specify functionality of logic at a high level
- ✓ **Logic synthesis**
 - process of compiling HDL program into logic gates and flip-flops
- ✓ **Validation** - high-level **simulation** to catch specification errors
 - verify pin-outs and connections to other system components
 - low-level to verify and check performance



CAD Tool Path (cont'd)


- ✓ Technology mapping
 - map the logic onto elements available in the implementation technology (LUTs for Xilinx FPGAs)
- ✓ Placement and routing
 - assign logic blocks to functions
 - make wiring connections
- ✓ Partitioning and constraining
 - if design does not fit or is unroutable as placed split into multiple chips
 - if design is too slow prioritize critical paths, fix placement of cells, etc.
 - few tools to help with these tasks exist today
- ✓ Generate programming files - bits to be loaded into chip for configuration



High-level synthesis



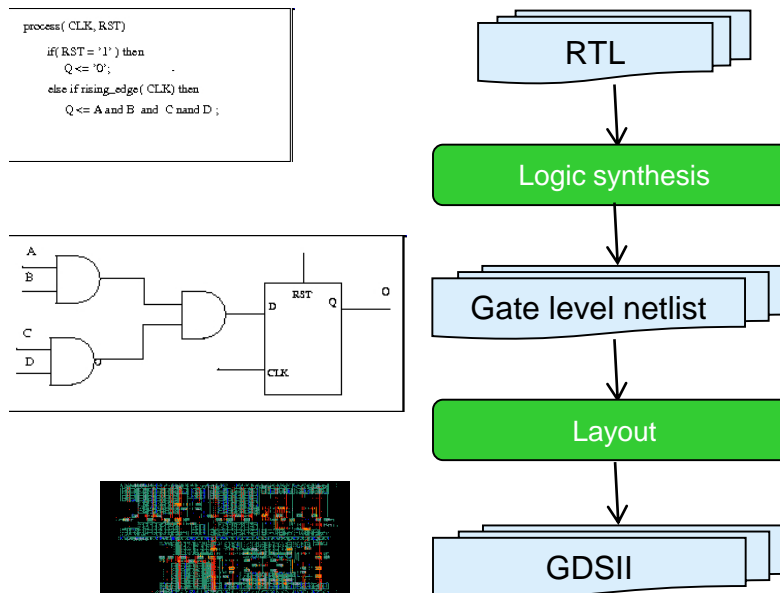
Electronic System Level Design (ESLD)





Typical HW design flow

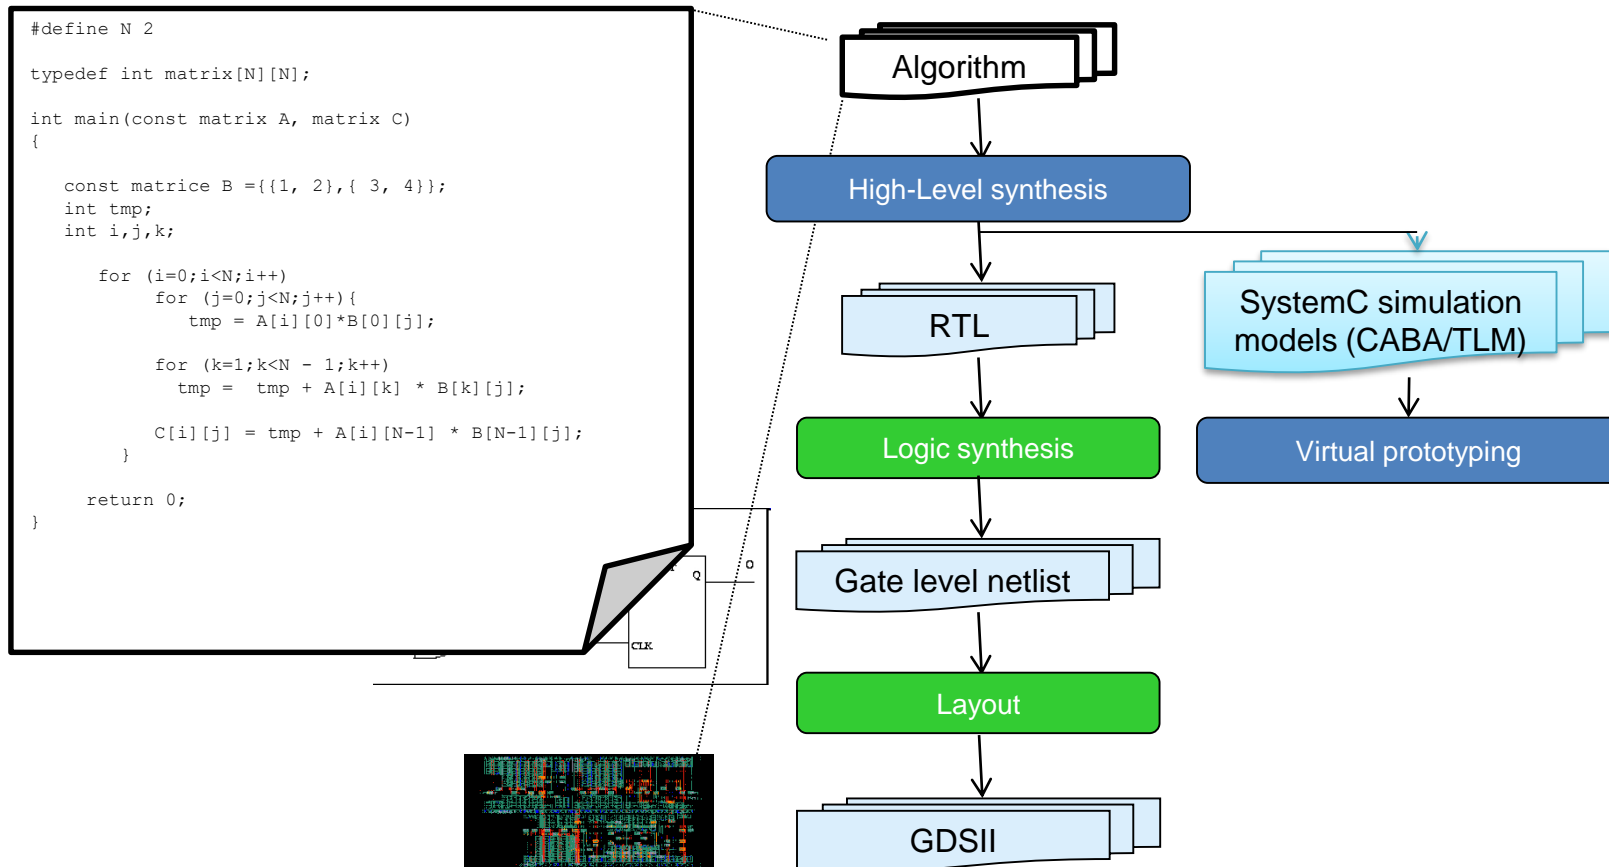
- ✓ Starting from a Register Transfer Level description, generate an IC layout





Typical HW design flow

- ✓ Starting from a Register Transfer Level description, generate an IC layout





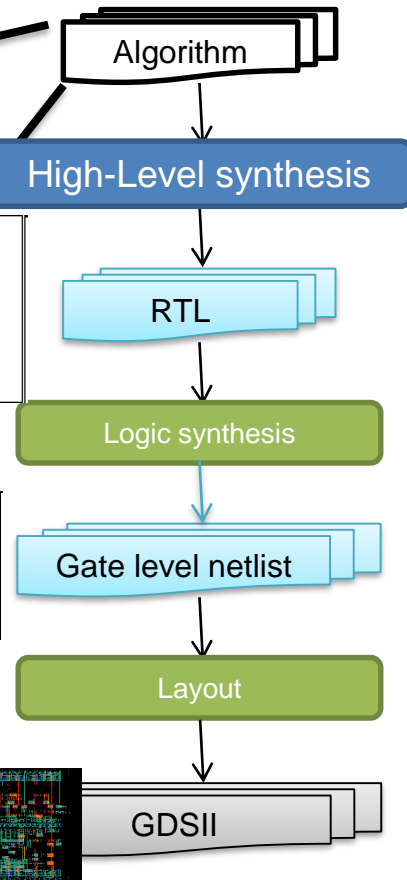

High-Level Synthesis

- ✓ Starting from C code, generate RTL, and then, layout
 - Typically, HW accelerators (FFT?) Ips
 - Can also synthesize a fully fledged processor!!
 - Soft-cores

```
#define N 2
typedef int matrix[N][N];

int main(const matrix A, matrix C)
{
    const matrix B = {{1, 2},{ 3, 4}};
    int tmp;
    int i,j,k;
    for (i=0;i<N;i++)
        for (j=0;j<N;j++) {
            tmp = A[i][0]*B[0][j];
            for (k=1;k<N - 1;k++)
                tmp = tmp + A[i][k] * B[k][j];
            C[i][j] = tmp + A[i][N-1] *
                B[N-1][j];
        }
    return 0;
}
```

```
process( CLK, RST)
if( RST = '1') then
    Q <= '0';
else if rising_edge( CLK) then
    Q <= A and B and C and D;
```





It is basically, a compiler!

- ✓ From C code
 - Generates the "physical" representation of Hardware modules
 - Registry Transfer Level, RTL
 - That will be deployed on the board

- ✓ Automatically



Synthesizable C subset

- ✓ No pointers
 - Statically unresolved
 - Arrays are allowed!
- ✓ No standard function call
 - printf, scanf, fopen, malloc...
- ✓ Function calls are allowed
 - Can be in-lined or not
- ✓ Nearly all datatypes are allowed
 - Specific datatypes are encouraged
 - Bit accurate integers, fixed point, signed, unsigned...



Example #1: a simple C code

```
#define N 16

int main(int data_in, int *data_out) {
    static const int Coeffs [N] = { 98, -39, -327, 439, 950, -2097, -1674, 9883,
                                     9883, -1674, -2097, 950, 439, -327, -39, 98 };

    int Values[N];
    int temp;
    int sample,i,j;

    sample = data_in;
    temp = sample * Coeffs[N-1];

    for(i = 1; i<=(N-1); i++) {
        temp += Values[i] * Coeffs[N-i-1];
    }

    for(j=(N-1); j>=2; j-=1 ) {
        Values[j] = Values[j-1];
    }

    Values[1] = sample;
    *data_out=temp;

    return 0;
}
```




Example #2: bit accurate C++ code

```
#include "ac_fixed.h" // From Mentor Graphics
#define PORT_SIZE ac_fixed<16, 12, true, AC_RND, AC_SAT> // 16 bits, 12 bits after the \
    point, quantization = rounding, overflow = saturation
#define N 16

int main(PORT_SIZE data_in, PORT_SIZE &data_out) {
    static const PORT_SIZE Coeffs [N]= { 1.1, 1.5, 1.0, 1.0, 1.7, 1.8, 1.2, 1.0,
                                           1.6, 1.0, 1.5, 1.1, 1.9, 1.3, 1.4, 1.7 };

    PORT_SIZE Values[N];
    PORT_SIZE temp;
    PORT_SIZE sample;

    sample= data_in;
    temp = sample * Coeffs[N-1];

    for(int i = 1; i<=(N-1); i++) {
        temp = Values [i] * Coeffs[N-i-1] + temp;
    }

    for(int j=(N-1); j>=2; j-=1 ) {
        Values[j] = Values [j-1];
    }

    Values[1] = sample;
    data_out=temp;

    return 0;
}
```



High-level transformations

✓ Loops

- Loop pipelining,
- Loop unrolling
- Loop merging
- Loop tiling
- ...

✓ Arrays mapping

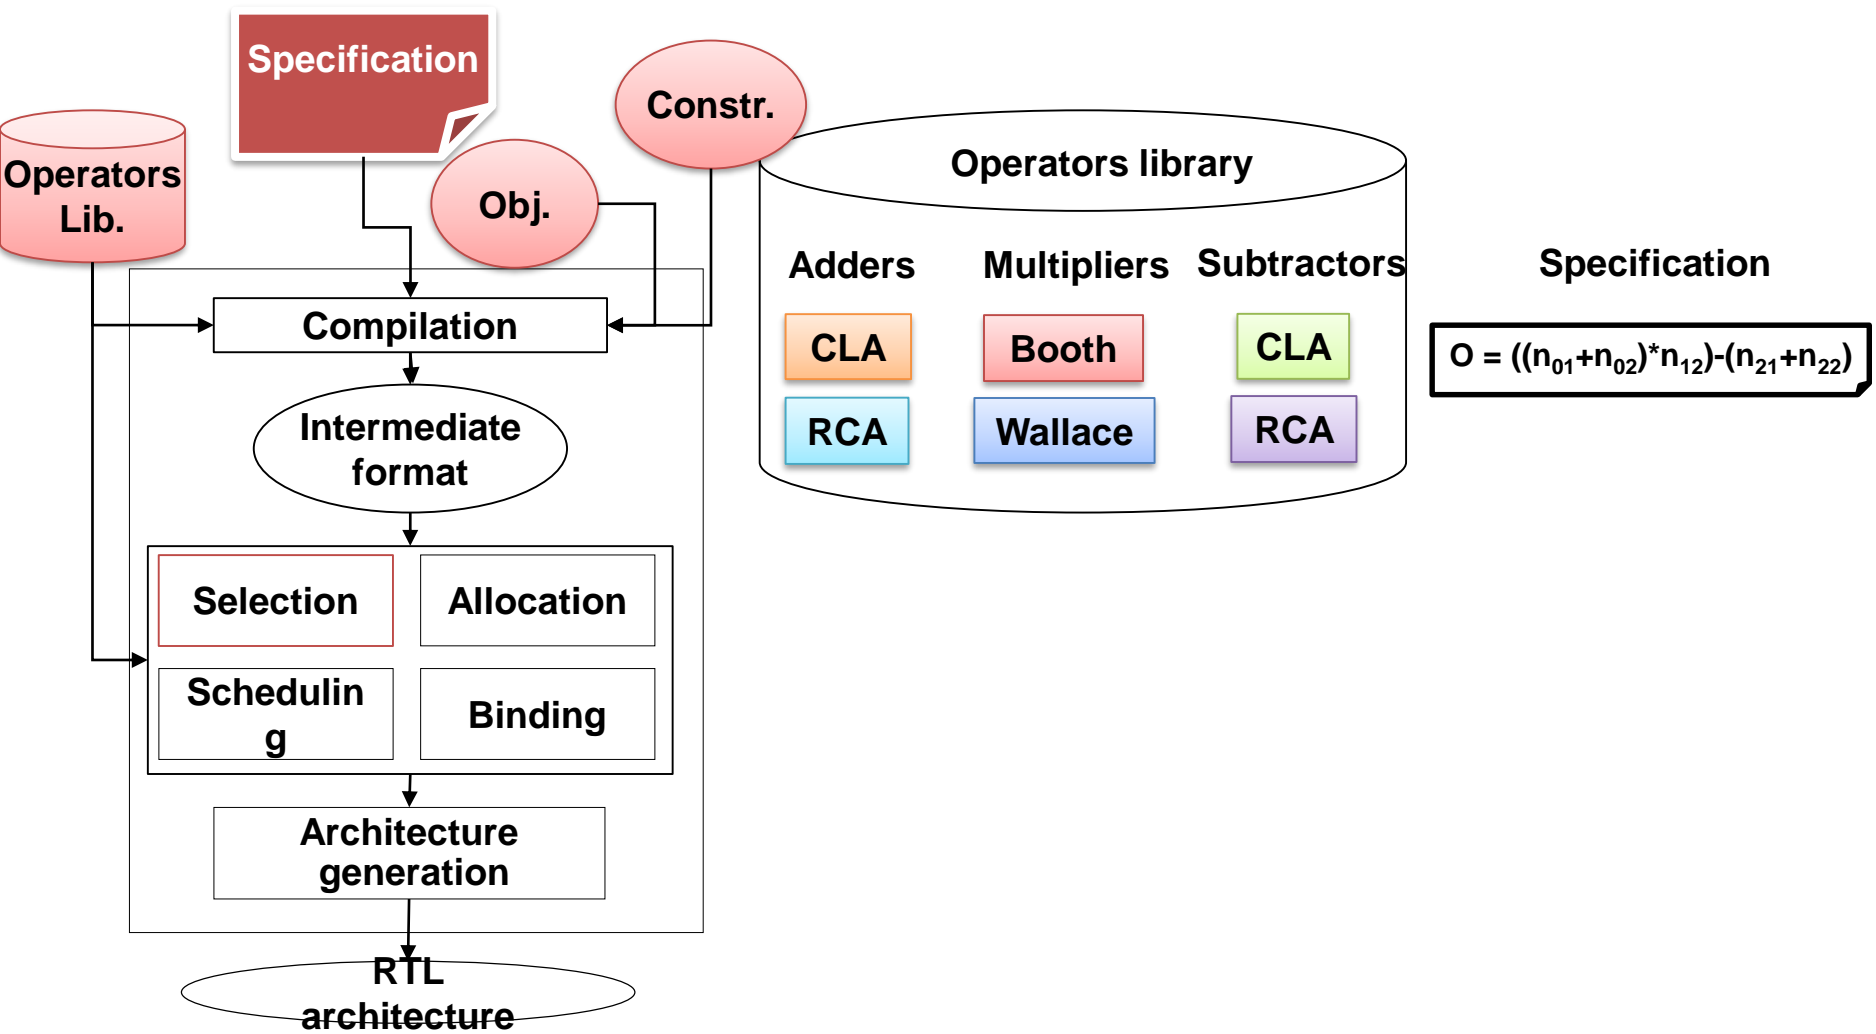
- Arrays can be mapped on memory banks
- Arrays can be synthesized as registers
- Constant arrays can be synthesized as logic
- ...

✓ Functions

- Function calls can be in-lined
- Function is synthesized as an operator
 - Sequential, pipelined, functional unit...
- Single function instantiation
- ...

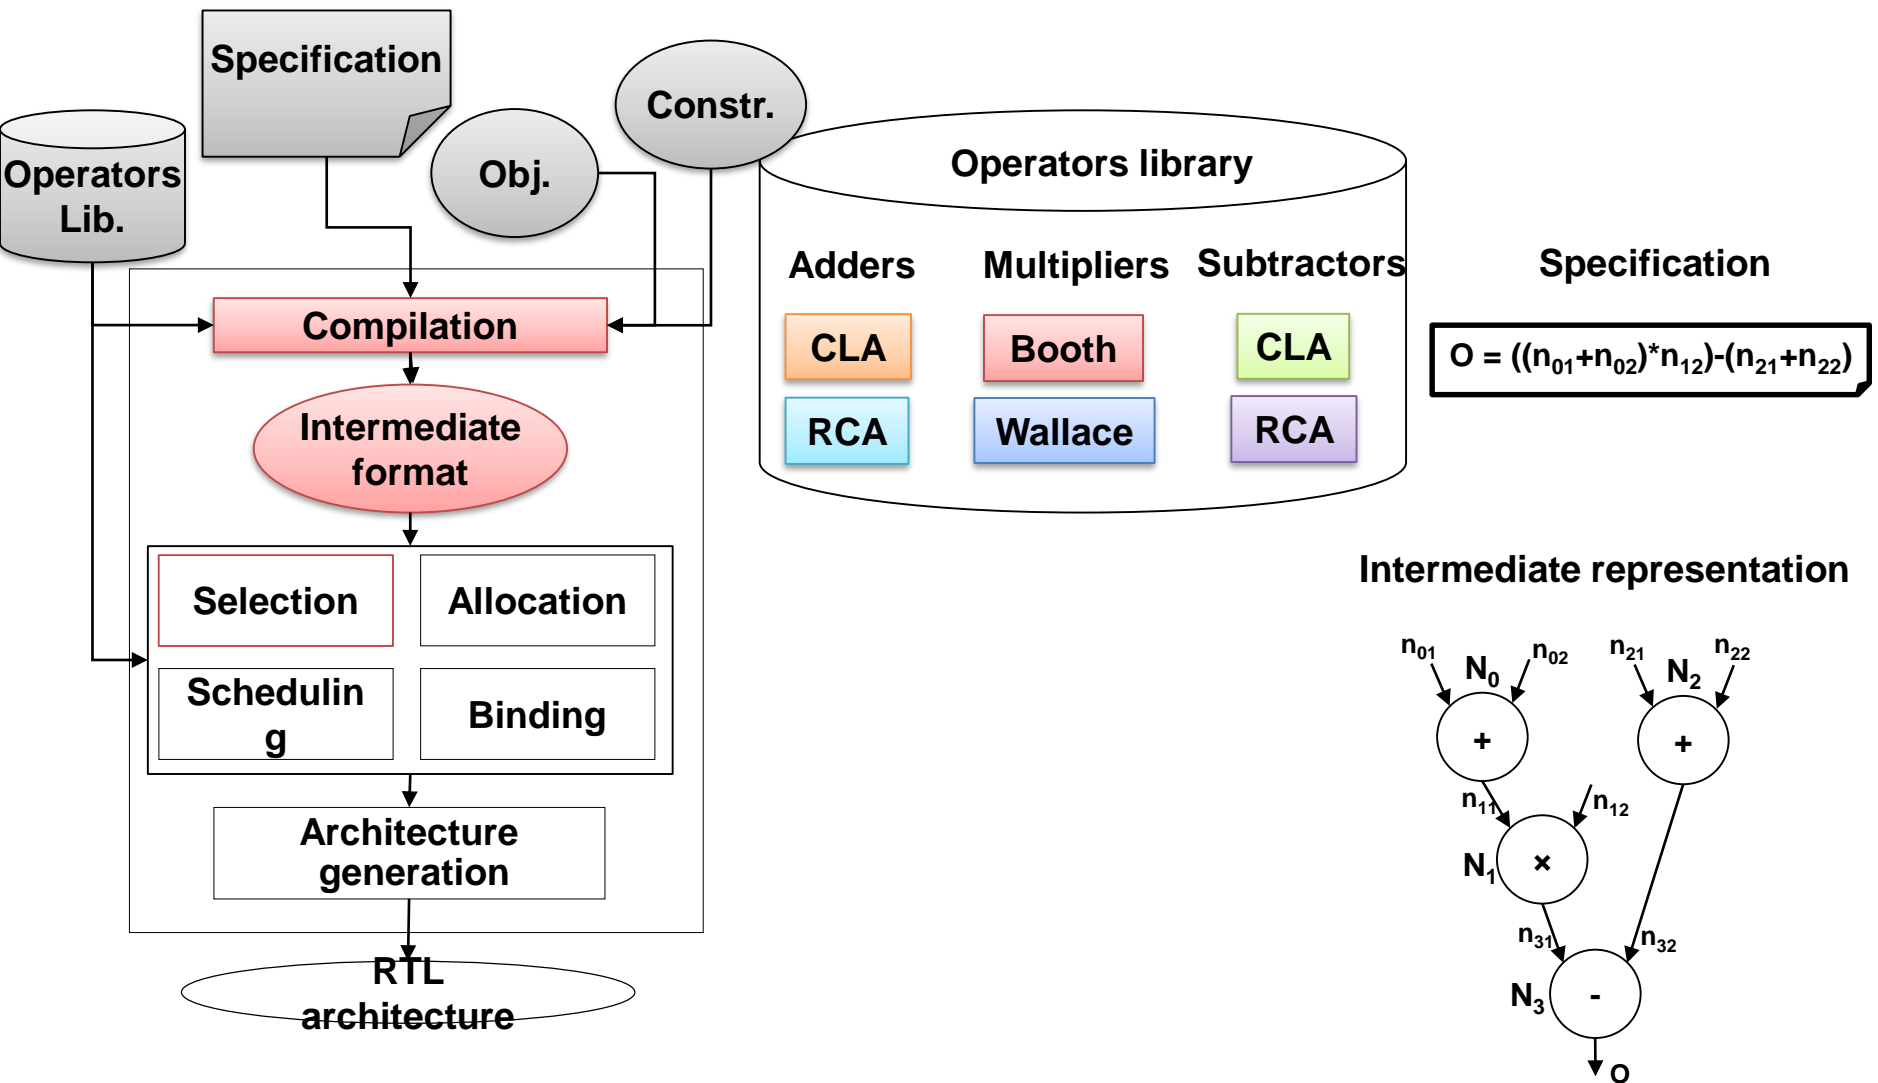


HLS steps: Inputs



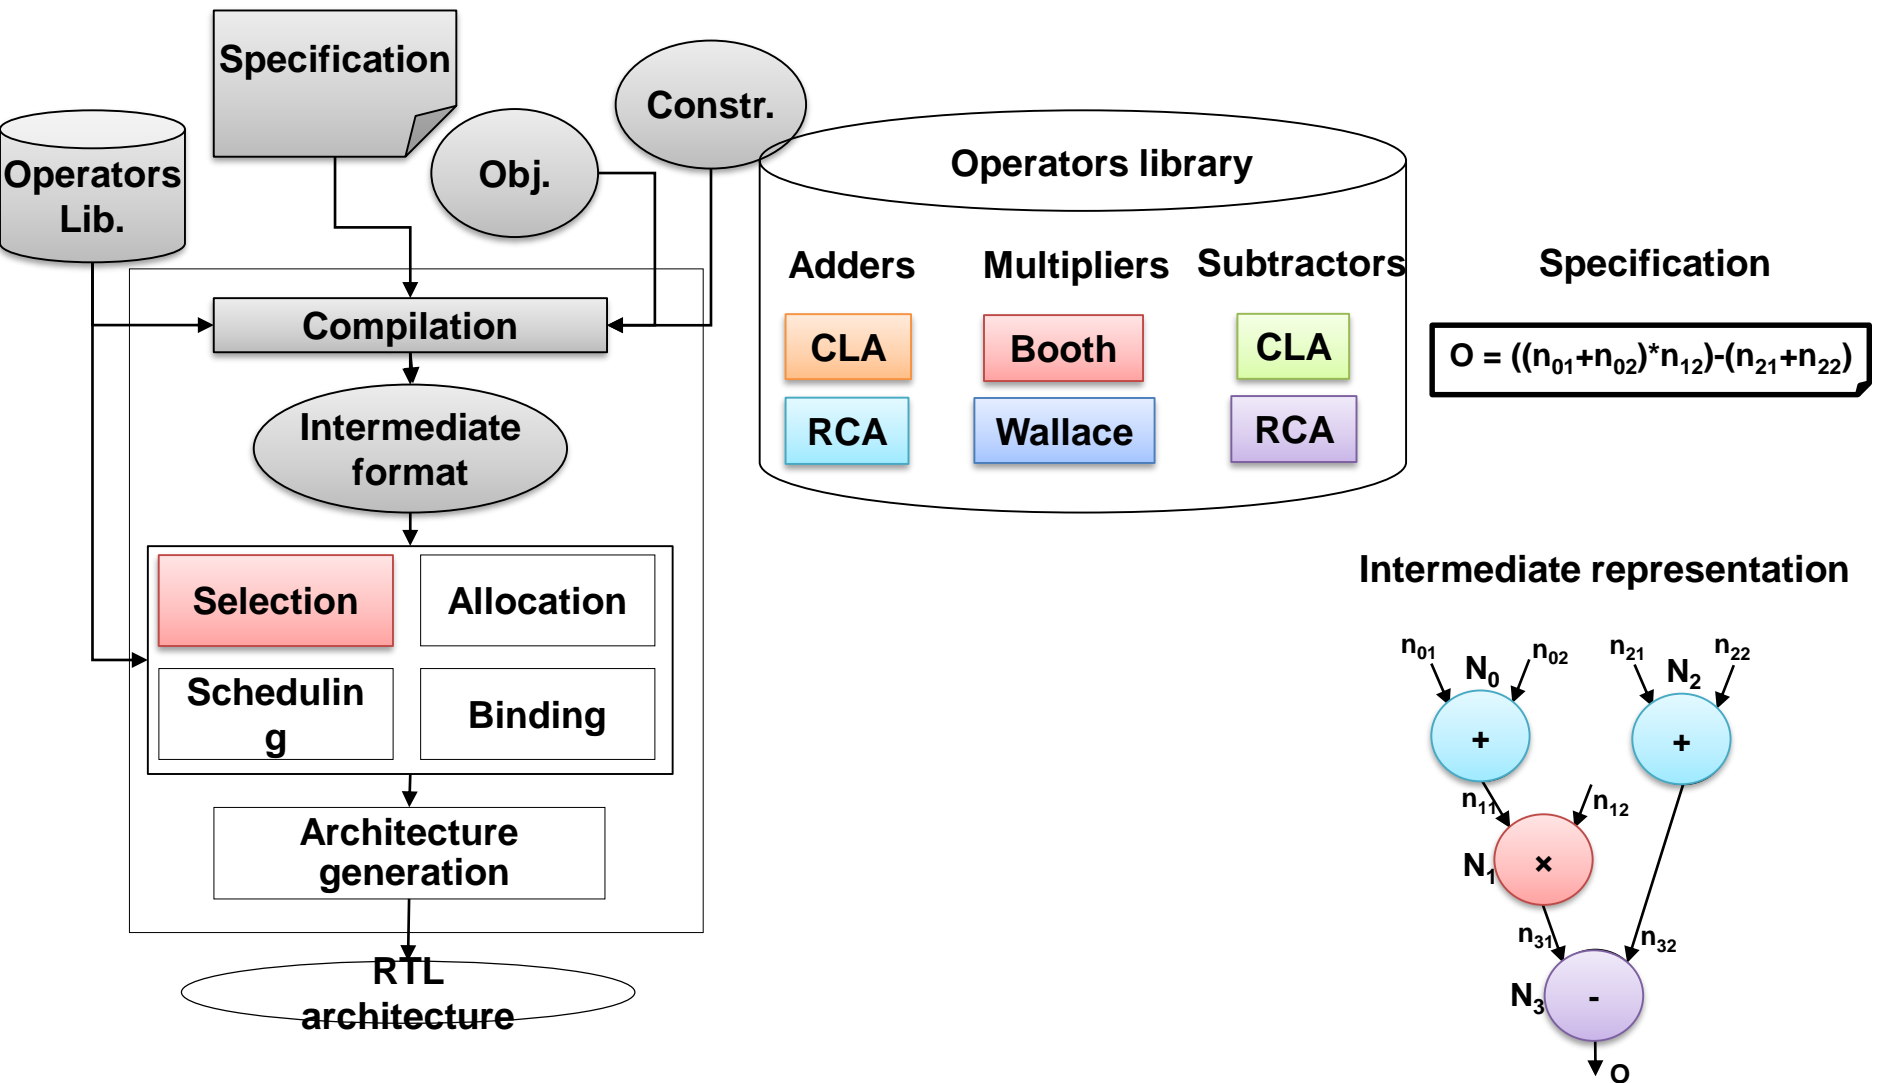


HLS steps: Compilation



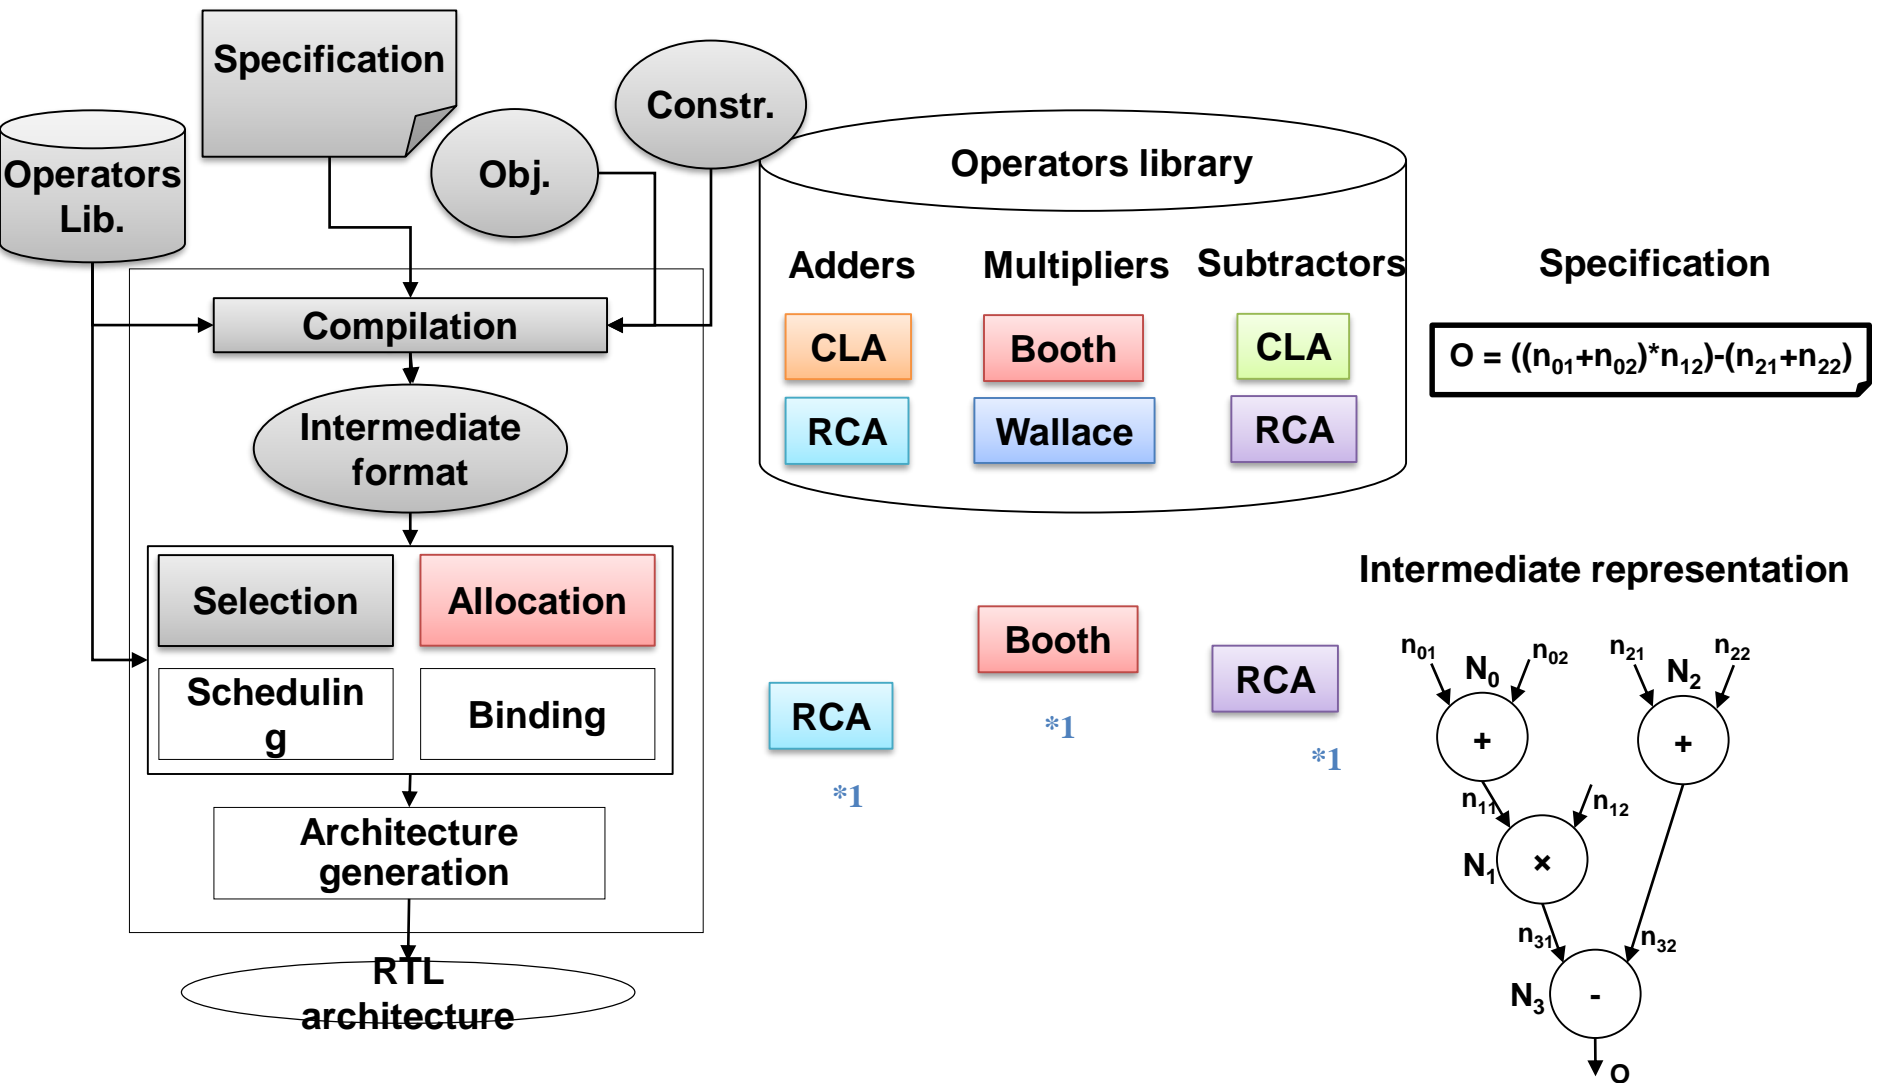


HLS steps: Selection



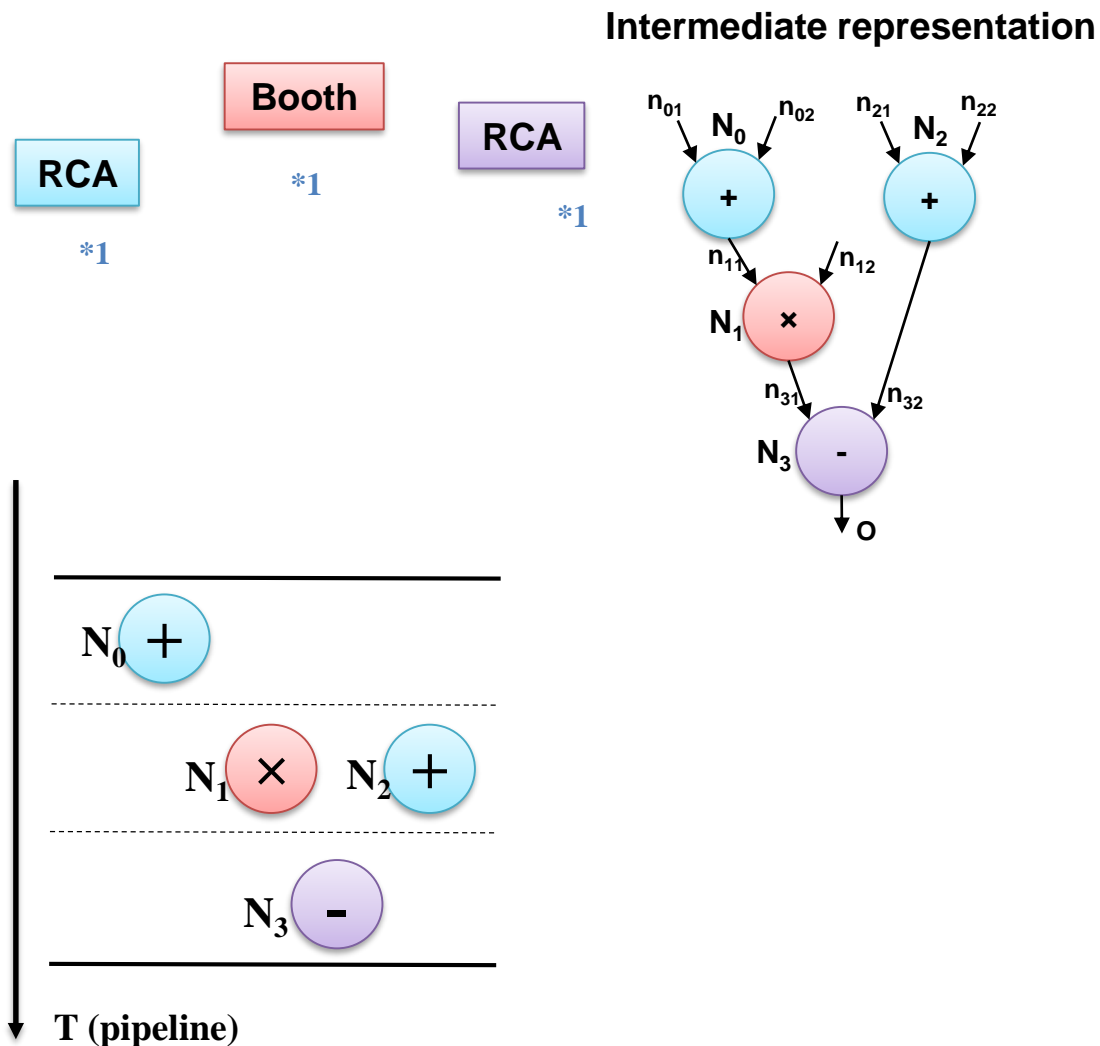
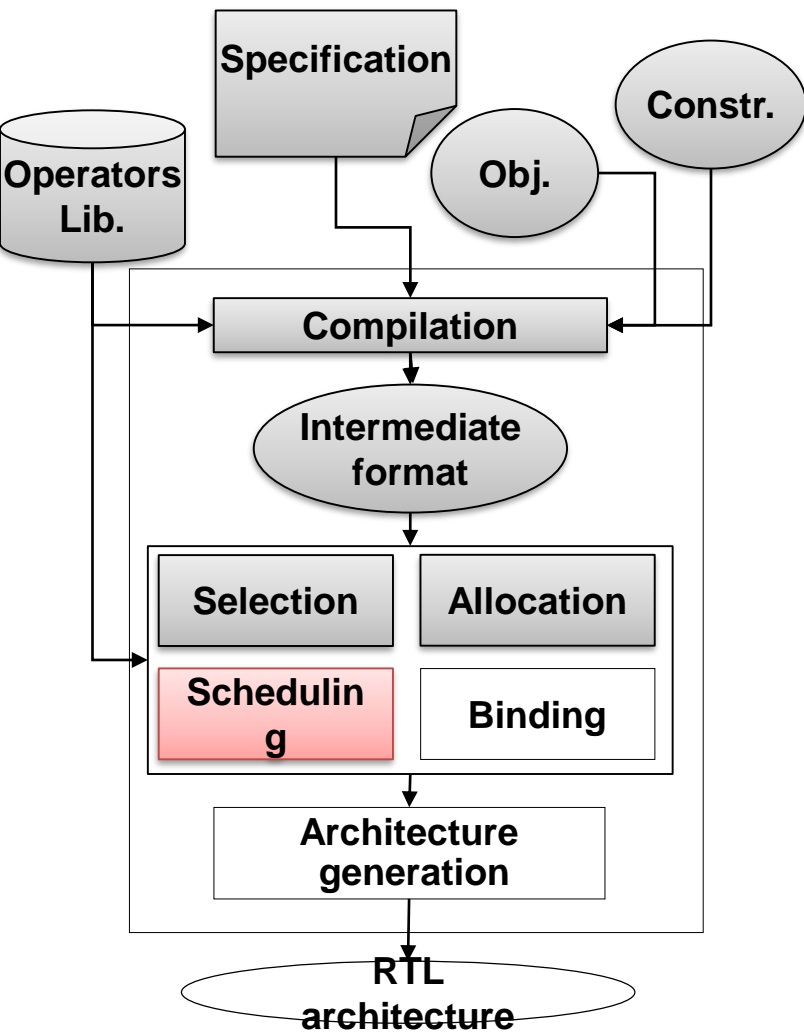


HLS steps: Allocation




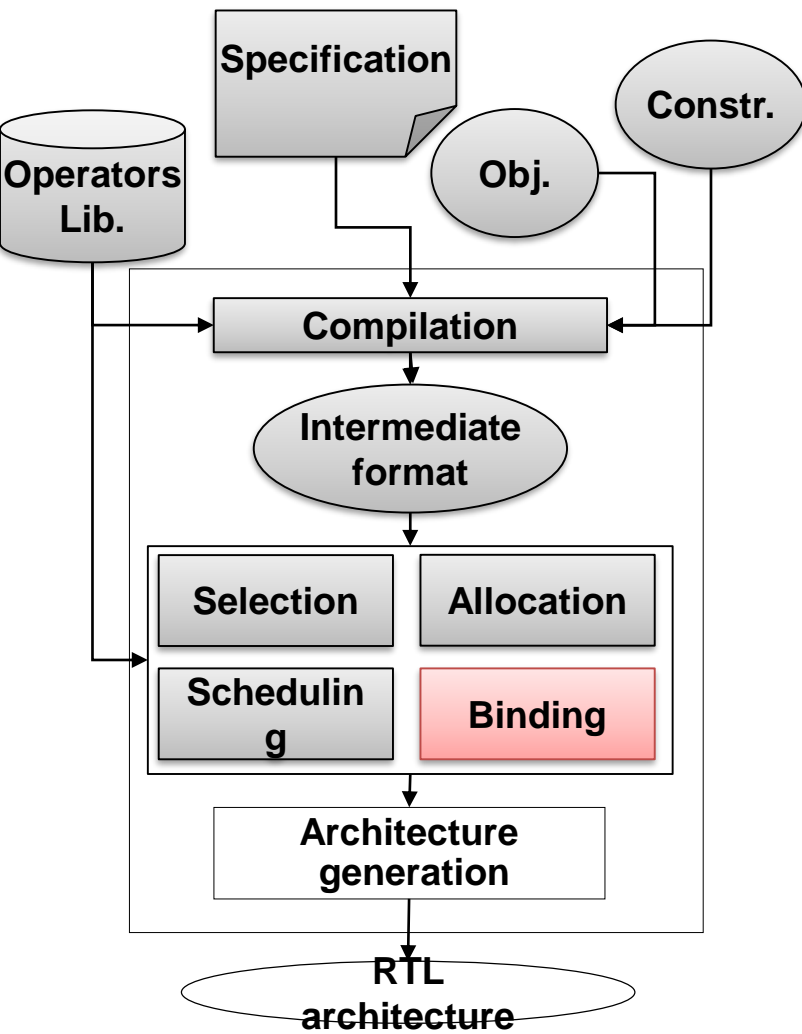


HLS steps: Scheduling






HLS steps: Binding

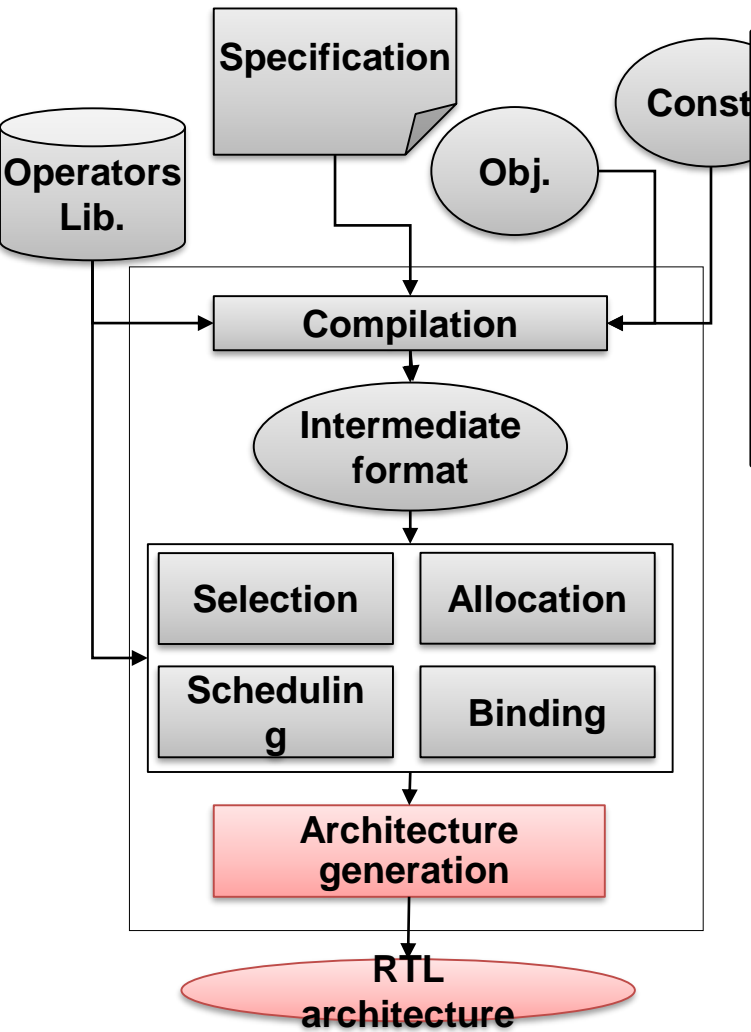


Intermediate representation




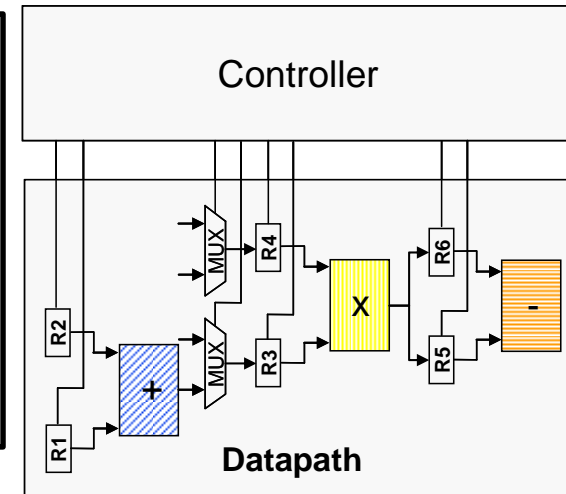


HLS steps: Binding



- Controller**
- **FSM controller**
 - **Programmable controller**

- Datapath components**
- **Storage components**
 - **Functional units**
 - **Connection components**



Data Binding

$n_{01} \rightarrow R_1$
 $n_{02} \rightarrow R_2$
 $n_{21}, n_{11} \rightarrow R_3$
 $n_{22}, n_{12} \rightarrow R_4$
 $n_{31} \rightarrow R_5$
 $n_{32} \rightarrow R_6$



Xilinx's Vivado SDK

The FPGA development tool

- ✓ Starting from C or RTL...
 - ✓ ...generates and deploys the IP on the FPGA
 - ✓ ..as well as SW artifacts to interact with them (drivers)
-
- ✓ Let's see it in action!




Heterogeneous systems



Host-accelerator model

- ✓ Multi-core General purpose host
 - The "traditional" core
- ✓ Coupled with a co-processor/accelerator







Something you are used to

GP-GPU based systems

- ✓ As in your laptop
 - ...yes, the one under your nose....
- ✓ Host => control-based code
- ✓ GPU => regular, highly-parallel code






Something you are less used to

GP-GPU based embedded platforms


- ✓ ...this is not under your nose....
- ✓ Still, host + accelerator model
- ✓ Communicate via shared memory
 - No PCI-express
 - Host memory "pull"
 - FPGA mem/BRAM "push"





FPGA-based accelerators

- ✓ Can create hundreds of small HW accelerators (de/crypt, de/coders)
- ✓ Can even create a single core (as co-processor)
 - Soft-cores
- ✓ Communicate via shared memory
 - No PCI-express
 - Host memory "pull"
 - FPGA mem/BRAM "push"





Example: Xilinx Zynq

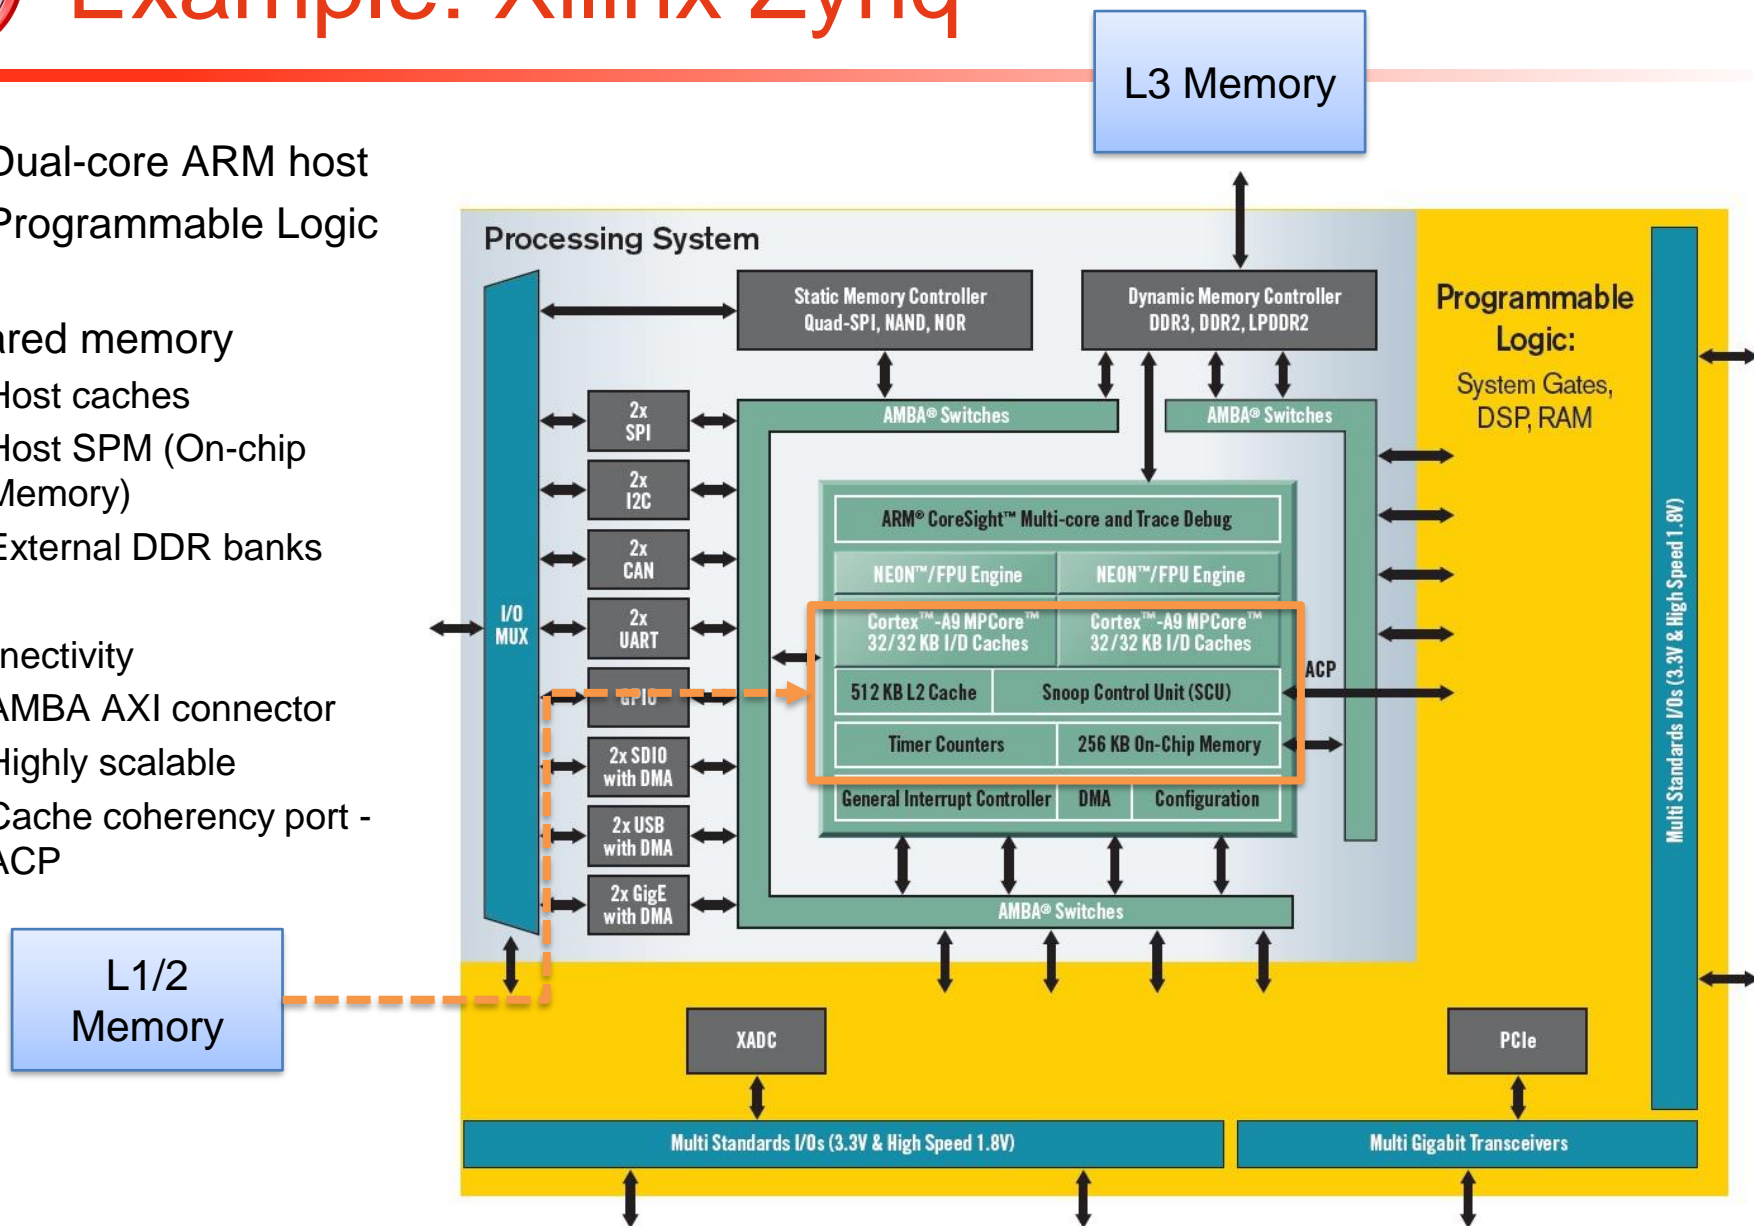
- ✓ Dual-core ARM host
- ✓ Programmable Logic

Shared memory

- ✓ Host caches
- ✓ Host SPM (On-chip Memory)
- ✓ External DDR banks

Connectivity

- ✓ AMBA AXI connector
- ✓ Highly scalable
- ✓ Cache coherency port - ACP






Xilinx FPGA SoCs



Xilinx Zynq-7000

- ✓ Dual-core ARM host
- ✓ Programmable Logic
- ✓ Shared memory






Xilinx Zynq-7000

FAMILY	PART	Logic Cells (K)	Block RAM (Mb)	DSP Slices	Maximum I/O Pins	Maximum Transceiver (Video Code Unit (VCU)
ZYNQ-7000							
	Z-7010	28	2,1	80	100	-	-
	Z-7015	74	3,3	160	150	4	-
	Z-7020	85	4,9	220	200	-	-
	Z-7030	125	9,3	400	250	4	-
	Z-7035	275	17,6	900	362	16	-
	Z-7045	350	19,1	900	362	16	-
	Z-7100	444	26,5	2020	400	16	-



Xilinx Zynq Ultrascale+





Xilinx Zynq Ultrascale portfolio

✓ Zynq UltraScale+ CG

- Dual-core Cortex-A53 and a dual-core Cortex-R5 real-time processor
- Programmable logic
- Optimized for industrial motor control, sensor fusion, and industrial IoT applications

✓ Zynq UltraScale+ EG

- Quad-core Cortex-A53 and dual-core Cortex-R5 real-time processors
- Mali-400 MP2 graphics processing unit + programmable logic
- Next-generation wired and 5G wireless infrastructure, cloud computing, and Aerospace and Defense applications

✓ Zynq UltraScale+ EV

- EG platform + integrated H.264 / H.265 video codec
- Multimedia, automotive ADAS, surveillance, and other embedded vision applications



Xilinx Zynq Ultrascale+

FAMILY	PART	Logic Cells (K)	Block RAM (Mb)	DSP Slices	Maximum I/O Pins	Maximum Transceiver (Gbps)	Video Code Unit (VCU)
ZYNQ UltraScale+ CG							
	ZU2CG	103	5,3	240	252 -	-	-
	ZU3CG	154	7,6	360	252 -	-	-
	ZU4CG	192	18,5	728	252 -	-	-
	ZU5CG	256	23,1	1248	252 -	-	-
	ZU6CG	469	25,1	1973	328 -	-	-
	ZU7CG	504	38	1728	464 -	-	-
	ZU9CG	600	32,1	2520	328 -	-	-
ZYNQ UltraScale+ EG							
	ZU2EG	103	5,3	240	252 -	-	-
	ZU3EG	154	7,6	360	252 -	-	-
	ZU4EG	192	18,5	728	252 -	-	-
	ZU5EG	256	23,1	1248	252 -	-	-
	ZU6EG	469	25,1	1973	328 -	-	-
	ZU7EG	504	38	1728	464 -	-	-
	ZU9EG	600	32,1	2520	328 -	-	-
	ZU11EG	653	43,6	2928	512 -	-	-
	ZU15EG	747	57,7	3528	328 -	-	-
	ZU17EG	926	56,7	1590	668 -	-	-
	ZU19EG	1143	70,6	1968	668 -	-	-
ZYNQ UltraScale+ EV							
	ZU4EV	192	18,5	728	252 -	-	1
	ZU5EV	256	23,1	1248	252 -	-	1
	ZU7EV	504	38	1728	464 -	-	1



Zedboard

- ✓ Complete development kit with Xilinx Zynq-7000 SoC
- ✓ Basic support for rapid prototyping and proof-of-concept development
- ✓ Small 😊





UltraZed

- ✓ System-On-Module (SOM)
 - ✓ Based on the Ultrascale architecture: no host!
 - ✓ Packages system memory, Ethernet, USB, and configuration memory needed for an embedded processing system
-
- ✓ UltraZed EG
 - ✓ UltraZed EV



Xilinx Pynq: Python for Zynq

- ✓ Open-source project from Xilinx for design
- ✓ Uses Python language and libraries
- ✓ Maximizes productivity

Processor: Dual-Core ARM® Cortex®-A9

FPGA: 1.3 M reconfigurable gates

Memory: 512MB DDR3 / FLASH

Storage: Micro SD card slot

Video: HDMI In and HDMI Out

Audio: Mic in, Line Out

Network: 10/100/1000 Ethernet


Expansion: USB Host connected to ARM PS

Interfaces: 1x Arduino Header, 2x Pmod (49 GPIO)

GPIO: 16 GPIO (65 in total with Arduino and Pmods)

Other I/O: 6x User LEDs, 4x Pushbuttons, 2x Switches

Dimensions: 3.44" x 4.81" (87mm x 122mm)





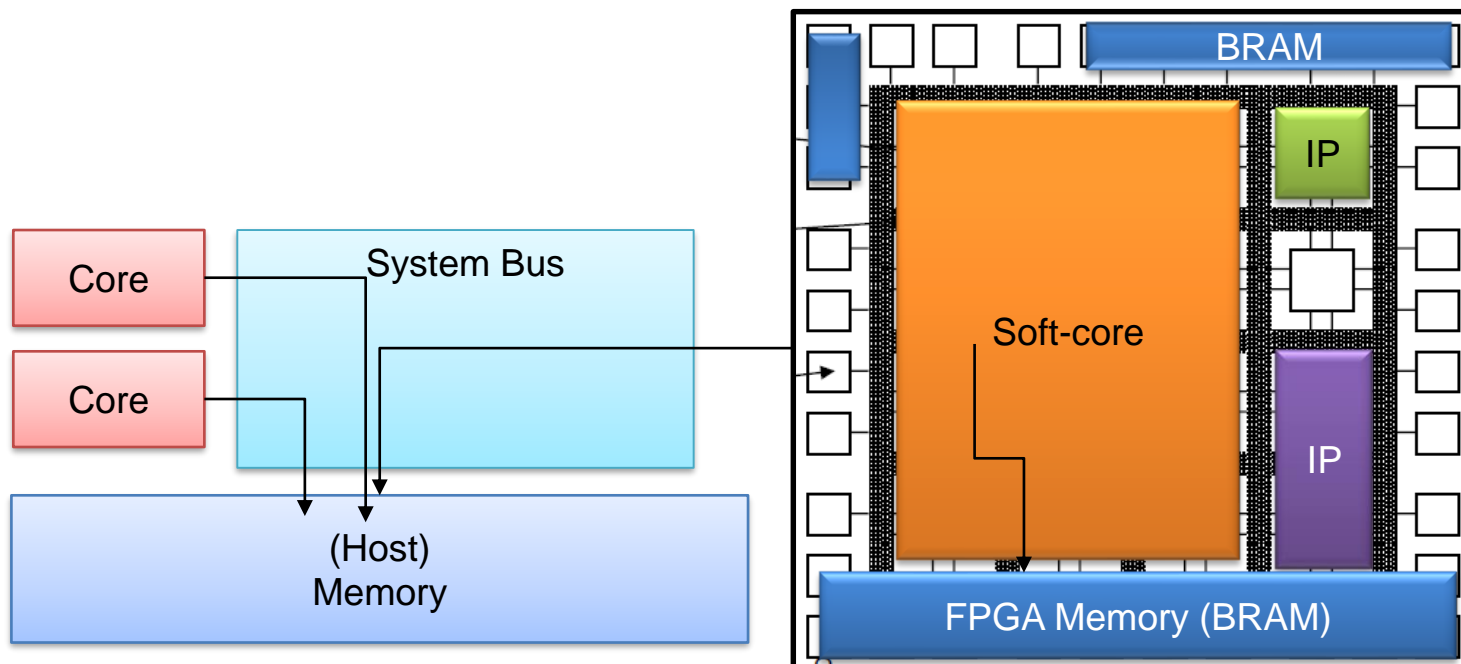
Programming heterogeneous systems



Heterogeneous programming

Besides a tool to generate the actual IPs, we need

- ✓ A way to efficiently offload (pre-compiled) bitcode on the FPGA
 - On-the-fly Dynamic Partial Reconfiguration (DPR)
- ✓ Simple offloading subroutines to the newly created HW blocks
 - To increase productivity
- ✓ In case we have SW cores, we need a toolchain to cross-compile for them





1) custom/"by hand"/CAD

Code generated by logic synthesis tool

- ✓ Step 1 – generate the bitcode of the accelerator
 - Vivado HLS
- ✓ Step 2 – plug the accelerator in a design
 - Vivado
 - Include processing system (ARM host) + accelerator + IC + ...
- ✓ Step 3 – generate the design
 - Bitcode ready to be installed of the IP
 - Architecture configuration files (memory maps...)
 - Software for host + drivers to communicate with the IP

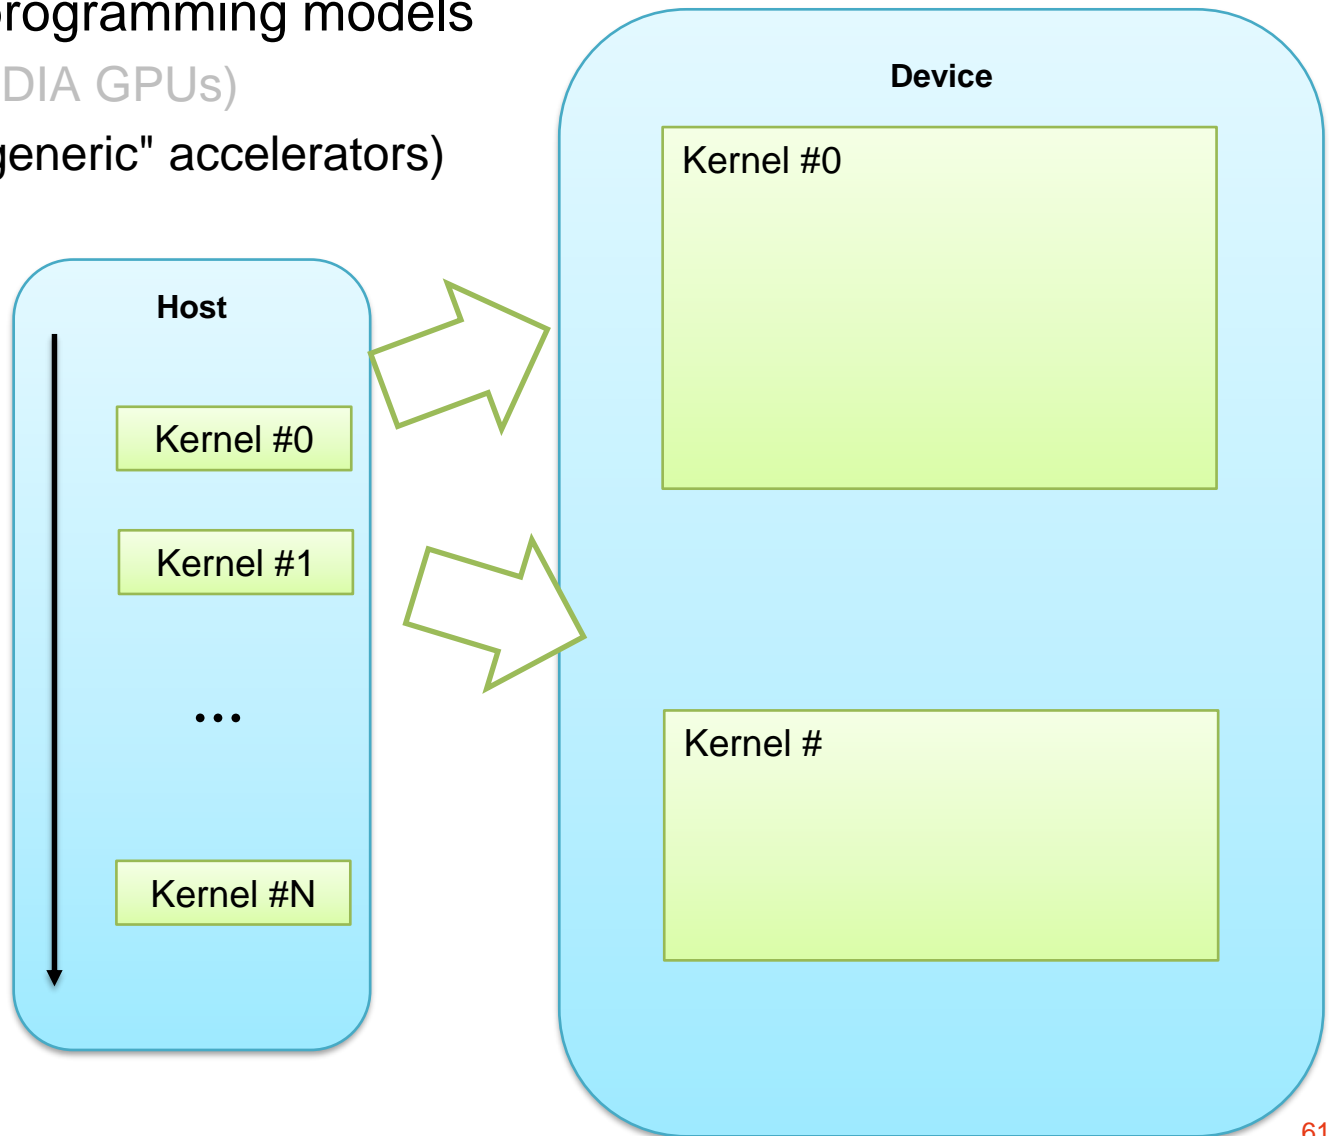
Let's see this in action!



2) offload-based programming

✓ Offload-based programming models

- CUDA (for NVIDIA GPUs)
- OpenCL (for "generic" accelerators)
- OpenMP 4.5





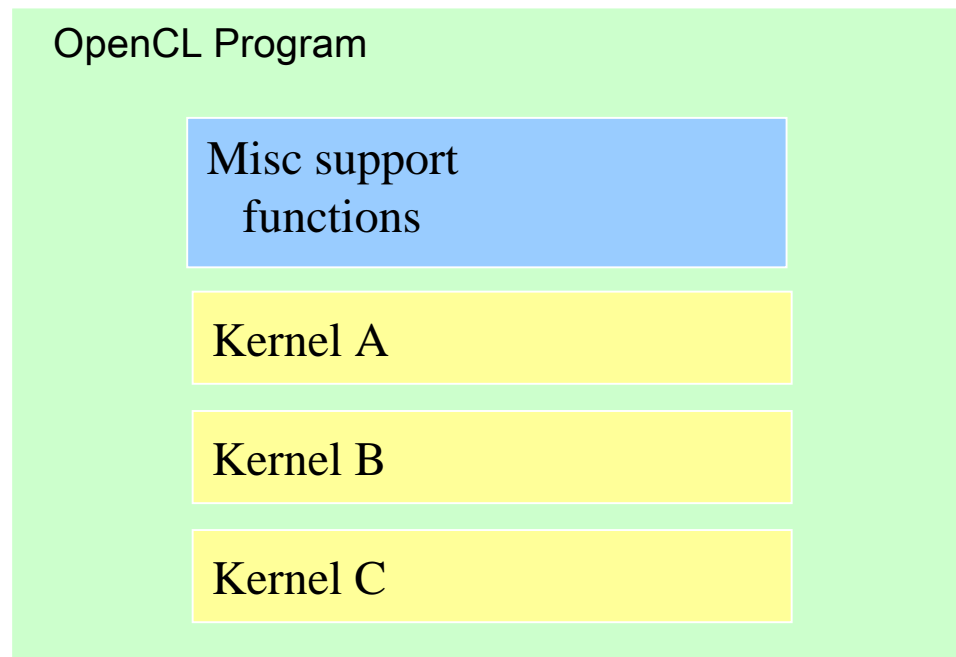
OpenCL

- ✓ OpenCL was initiated by Apple and maintained by the Khronos Group (also home of OpenGL) as an industry standard API
 - For cross-platform parallel programming in CPUs, GPUs, DSPs, FPGAs,...
- ✓ OpenCL host code is much more complex and tedious due to desire to maximize portability and to minimize burden on vendors



OpenCL program


- ✓ An OpenCL “program” is a C program that contains one or more “**kernels**” and any supporting routines that run on a target device
- ✓ An OpenCL kernel is the basic unit of parallel code that can be executed on a target **device**
- ✓ In our case, an FPGA





OpenCL execution model

- ✓ Integrated host+device app C program
 - Serial or modestly parallel parts in host C code
 - Highly parallel parts in device SPMD kernel C code
- ✓ Queues of command/data transfer to be executed on the device





OpenCL kernels – software version

- ✓ Code that executes on target devices
- ✓ Kernel body is instantiated N times (data parallel) – **work items**
- ✓ Each OpenCL work item gets a unique index
- ✓ In the FPGA case, we use IP drivers instead of this

```
__kernel void vadd(__global const float *a,  
                  __global const float *b,  
                  __global float *result)  
{  
    int id = get_global_id(0);  
    result[id] = a[id] + b[id];  
}
```



Host code – create exec ctx

```
cl_int clerr = CL_SUCCESS;
cl_context clctx = clCreateContextFromType(0, CL_DEVICE_TYPE_ALL, NULL, NULL,
&clerr);

size_t parmsz;
clerr = clGetContextInfo(clctx, CL_CONTEXT_DEVICES, 0, NULL, &parmsz);

cl_device_id* cldevs = (cl_device_id *) malloc(parmsz);
clerr = clGetContextInfo(clctx, CL_CONTEXT_DEVICES, parmsz, cldevs, NULL);

cl_command_queue clcmdq = clCreateCommandQueue(clctx, cldevs[0], 0, &clerr);
```



Host code – create data buffers

```
float *h_A = ..., *h_B = ...;
// allocate device (GPU) memory
cl_mem d_A, d_B, d_C;
d_A = clCreateBuffer(clctx, CL_MEM_READ_ONLY |
    CL_MEM_COPY_HOST_PTR, N *sizeof(float), h_A, NULL);
d_B = clCreateBuffer(clctx, CL_MEM_READ_ONLY |
    CL_MEM_COPY_HOST_PTR, N *sizeof(float), h_B, NULL);
d_C = clCreateBuffer(clctx, CL_MEM_WRITE_ONLY,
    N *sizeof(float), NULL, NULL);
```



Host code – device config setting

```
clkern=clCreateKernel(clpgm, "vadd", NULL);  
...  
clerr= clSetKernelArg(clkern, 0, sizeof(cl_mem), (void *)&d_A);  
clerr= clSetKernelArg(clkern, 1, sizeof(cl_mem), (void *)&d_B);  
clerr= clSetKernelArg(clkern, 2, sizeof(cl_mem), (void *)&d_C);  
clerr= clSetKernelArg(clkern, 3, sizeof(int), &N);
```



Host code – kernel launch

```
cl_event event=NULL;
clerr= clEnqueueNDRangeKernel(clcmdq, clkern, 2, NULL,
    Gsz, Bsz, 0, NULL, &event);
clerr= clWaitForEvents(1, &event);
clEnqueueReadBuffer(clcmdq, d_C, CL_TRUE, 0,
    N*sizeof(float), h_C, 0, NULL, NULL);
clReleaseMemObject(d_A);
clReleaseMemObject(d_B);
clReleaseMemObject(d_C);
}
```



OpenMP 4.5 (yr 2011)

```
#pragma omp target [clause [[,]clause]...] new-line  
    structured-block
```

Where clauses can be:

```
if([ target :] scalar-expression)  
device(integer-expression)  
private(list)  
firstprivate(list)  
map([[map-type-modifier[,]] map-type: ] list)  
is_device_ptr(list)  
defaultmap(tofrom:scalar)  
nowait  
depend(dependence-type: list)
```

- ✓ Introduces the concept of **device**
 - Execute structured block onto device
 - map clause to move data to-from the device
 - nowait for asynch execution



OpenMP 4.5 in action

- ✓ ESA application for infrared signal processing
 - Here, runs on a Kalray MPPA manycore

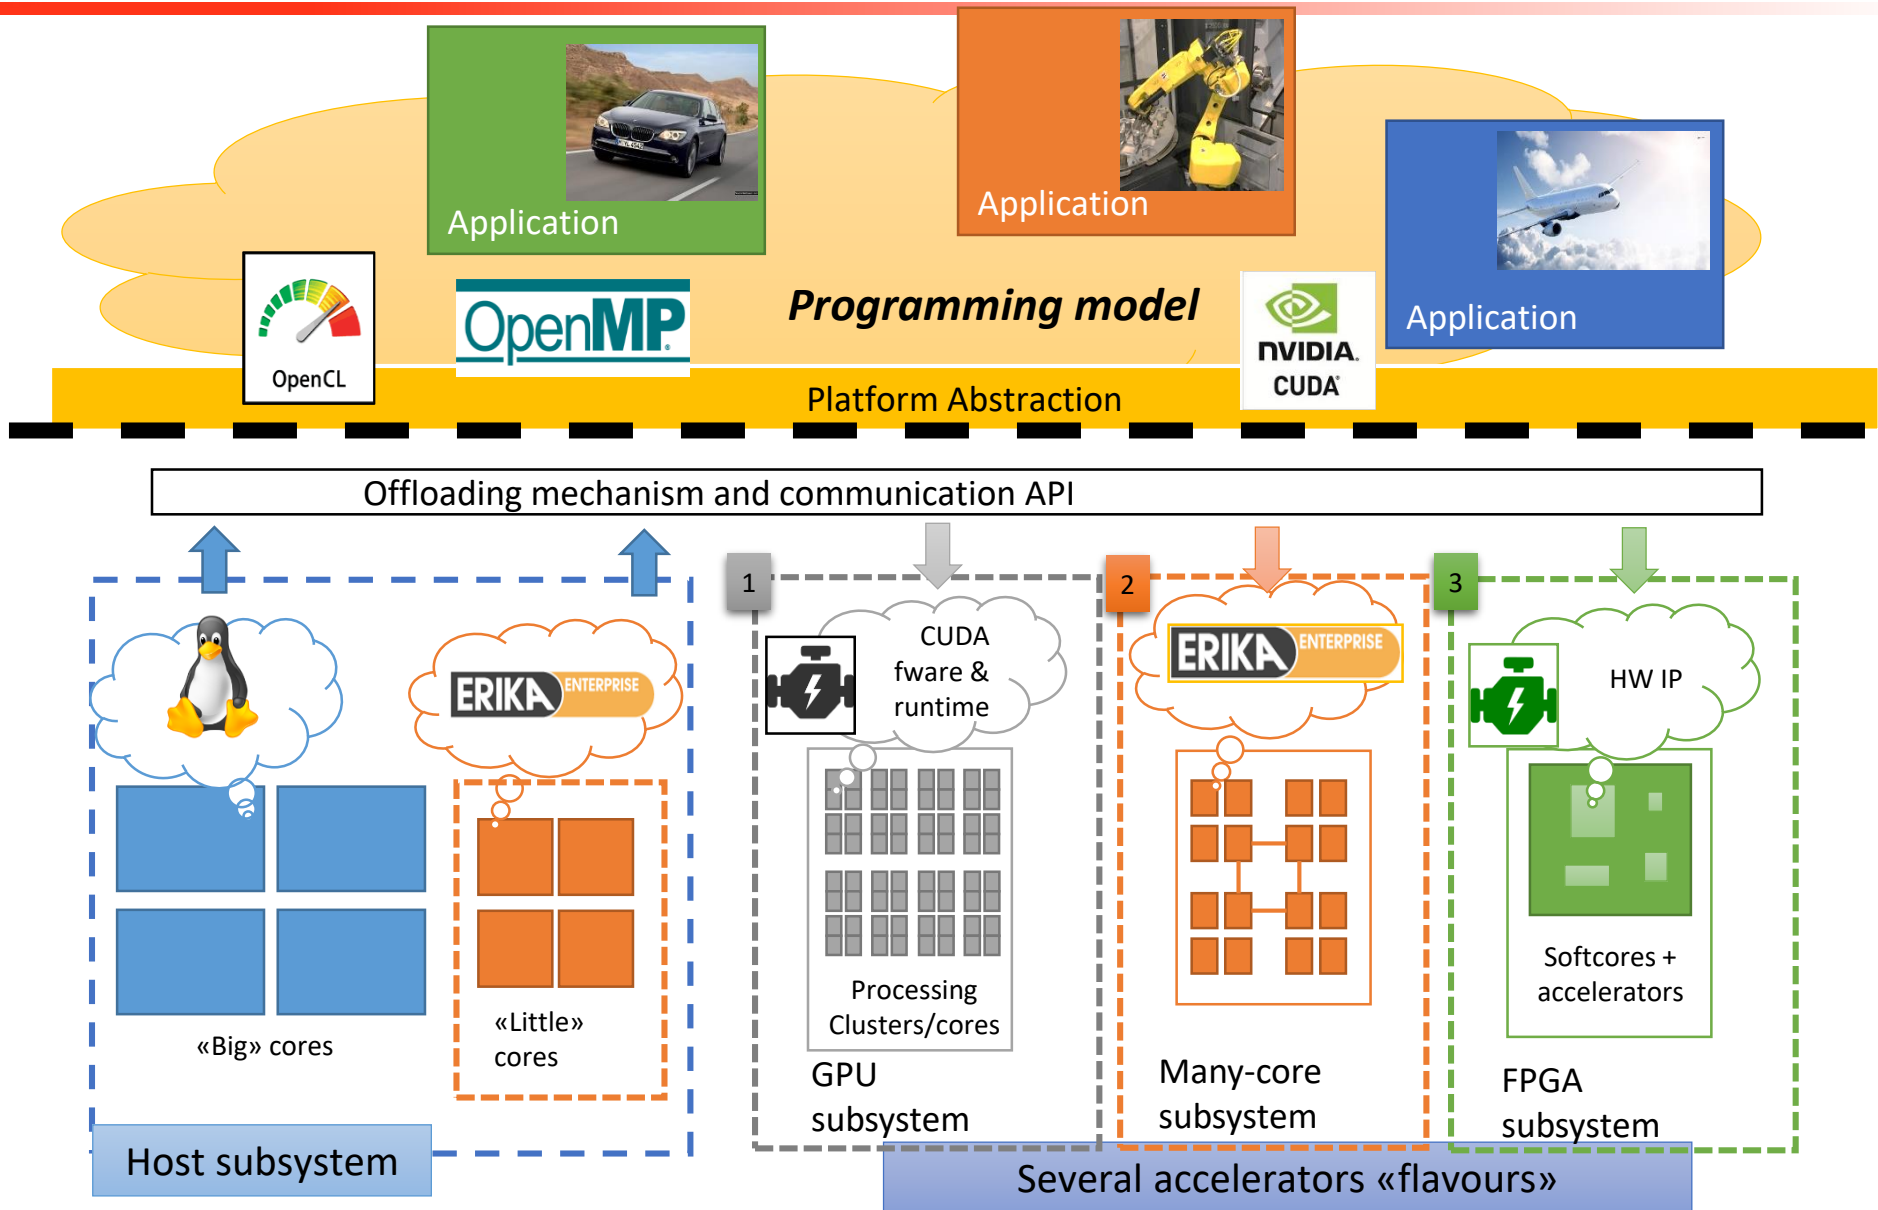
```
for (i = 0; i < DIM_Y; i++)
{
    for (j = 0; j < DIM_X; j=j+4)
    {
        UINT16BIT (*p_currentFrame1) [BS] = currentFrame[i][j];
        UINT16BIT (*p_currentFrame2) [BS] = currentFrame[i][j+1]

        #pragma omp target firstprivate(j) \
            map(to: saturationLimit[0:32]) \
            map(to: coeffOfNonLinearityPolynomial[0:32][0:4]) \
            map(tofrom: p_currentFrame1[0:bs][0:bs]) \
            map(tofrom: p_currentFrame2[0:bs][0:bs] \
            device(device_id) priority_id(0) nowait
        {
            phase1 (p_currentFrame1,p_currentFrame2
                saturationLimit, coeffOfNonLinearityPolynomial, j);
        }
    }
}
```



The Hercules framework

HERCULES



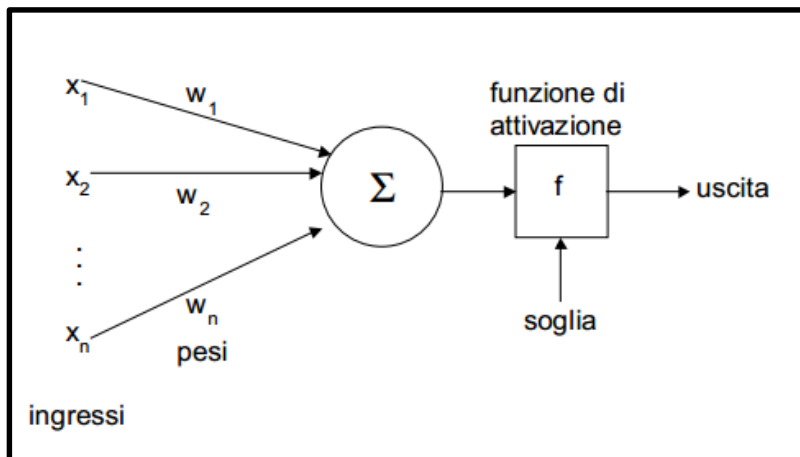



Neural Networks on FPGA accelerators



Neural networks

- ✓ Bio-inspired
- ✓ Based on neurons arranged in layers
 - And sub-layers
- ✓ Convolutional neural network
 - Neurons perform Convolutions






Convolution

- ✓ Computation-intensive
- ✓ Suitable for implementation in hardware
- ✓ In computer vision, blurring

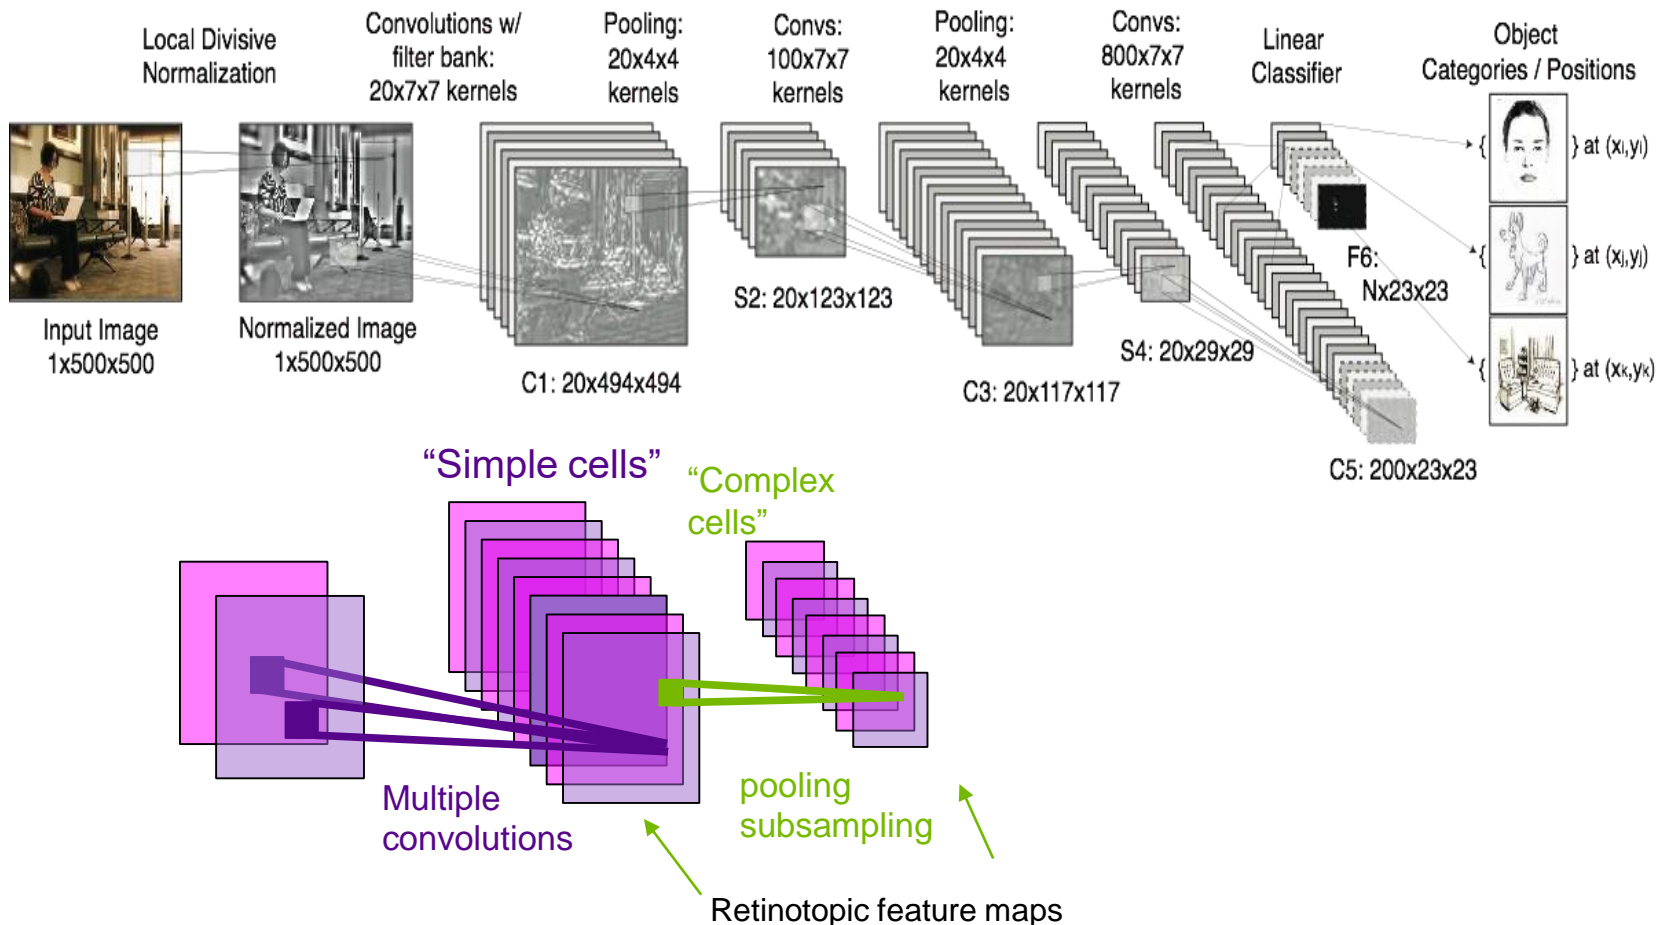
$$\begin{aligned}(f * g)(t) &\stackrel{\text{def}}{=} \int_{-\infty}^{\infty} f(\tau)g(t - \tau) d\tau \\ &= \int_{-\infty}^{\infty} f(t - \tau)g(\tau) d\tau.\end{aligned}$$








The convolutional net model

✓ (Multistage Hubel-Wiesel system)





The convolutional net model (cont'd)






Network parameters

✓ Network topology

- How many layers and sublayers?
- How big they are?
- How are they connected?

✓ Neuron type

- CNN
- int/float datatypes
- How to perform pooling?





Training a network

"The training problem"

- ✓ To set the weights/CNN kernels
- ✓ Training set must be huge

A "big data" problem

- ✓ Why do you think Google does self-driving cars?
- ✓ Why do you think big cloud players want our data?





CNN on FPGA

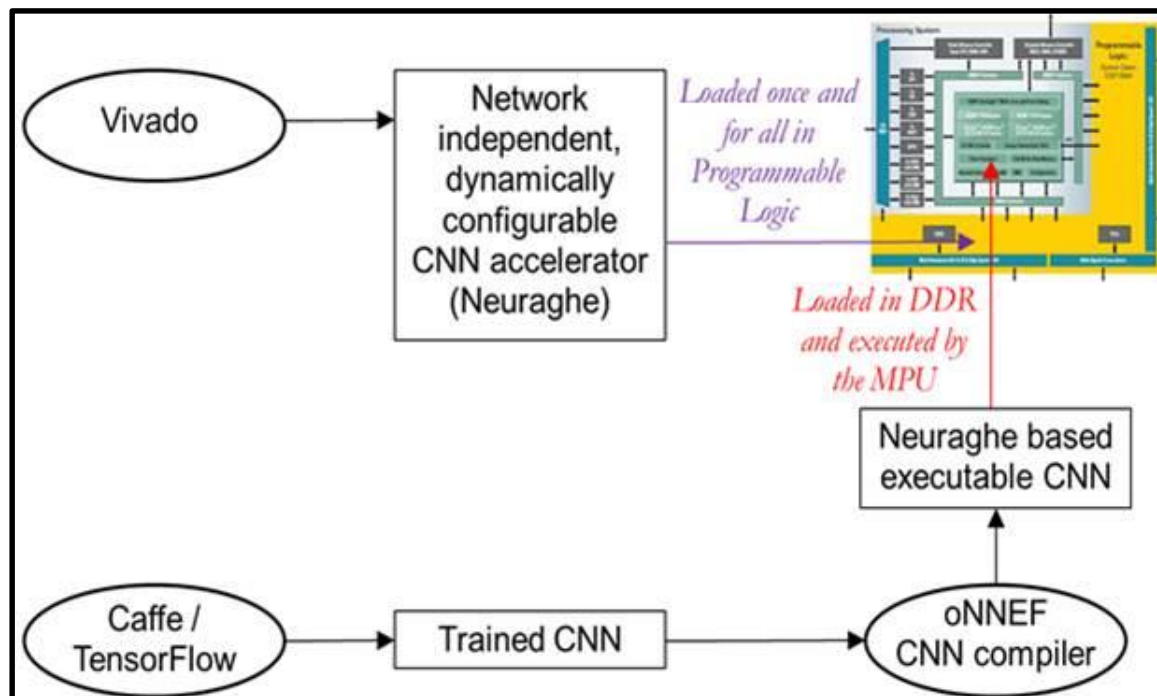
- ✓ We can implement one or more complete CNN layers on FPGA
 - How many?
- ✓ We can use float, int, datatypes
 - Int are smaller, but still efficient
- ✓ Binaries NN, where input weights are ± 1
 - Smaller, more suitable for area-constrained

Name	Supported Architecture
ZynqNet	XC-7Z045
BNN	Z-7020
GraphGen-based CNN	
Pre-Trained CNN based on LeNet5	
GoogLeNet	ZCU102
AlexNet	
VGG-16	
SSD-300	
FCN-AlexNet	



Nuraghe NN on FPGA

- ✓ Trained with ML frameworks Caffè and TensorFlow
- ✓ Nuraghe accelerator, configurable for specific CNN
- ✓ CNN compiler, which translates the CNN description from Caffè or TF in a program which runs on Nuraghe/Zynq





References



- ✓ "Calcolo parallelo" website
 - http://hipert.unimore.it/people/paolob/pub/Calcolo_Parallelo/
- ✓ My contacts
 - paolo.burgio@unimore.it
 - <http://hipert.mat.unimore.it/people/paolob/>
- ✓ Xilinx Zynq-7000 All Programmable SoC
 - <https://www.xilinx.com/products/silicon-devices/soc/zynq-7000.html>
- ✓ Pynq
 - <http://www.pynq.io/>
- ✓ Xilinx Ultrascale
 - <https://www.xilinx.com/products/technology/ultrascale.html>