# Barriers in OpenMP

Paolo Burgio paolo.burgio@unimore.it





## **Outline**

- > Expressing parallelism
  - Understanding parallel threads



- > Synchronization
  - Barriers, locks, critical sections
- > Work partitioning
  - Loops, sections, single work, tasks...
- > Execution devices
  - Target



## **OpenMP synchronization**

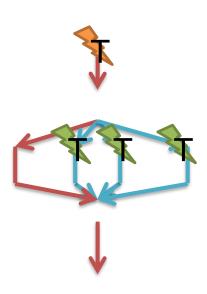
- > OpenMP provides the following synchronization constructs:
  - barrier
  - flush
  - master
  - critical
  - atomic
  - taskwait
  - taskgroup
  - ordered
  - ..and OpenMP locks



# Creating a parreg

- > Master-slave, fork-join execution model
  - Master thread spawns a team of Slave threads
  - They all perform computation in parallel
  - At the end of the parallel region, <u>implicit barrier</u>

```
int main()
{
    /* Sequential code */
    #pragma omp parallel num_threads(4)
    {
        /* Parallel code */
        // Parreg end: (implicit) barrier
        /* (More) sequential code */
}
```





## **OpenMP explicit barriers**

#pragma omp barrier new-line

(a standalone directive)

- > All threads in a team must wait for all the other threads before going on
  - "Each barrier region must be encountered by all threads in a team or by none at all"
  - "The sequence of barrier regions encountered must be the same for every thread in a team"
  - Why?
- > Binding set is the team of threads from the innermost enclosing parreg
  - "It applies to"
- > Also, it enforces a consistent view of the shared memory
  - We'll see this...



### **Exercise**

Let's code!

- > Spawn a team of (many) parallel Threads
  - Printing "Hello World"
  - Puta #pragma omp barrier
  - Reprint "Hello World" after
- > What do you see?
  - Now, remove the barrier construct
- > Now, put the barrier inside an if
  - E.g., if (omp\_get\_thread\_num() == 0) { ... }
  - What do you see?
  - Error!!!!



## **Effects on memory**

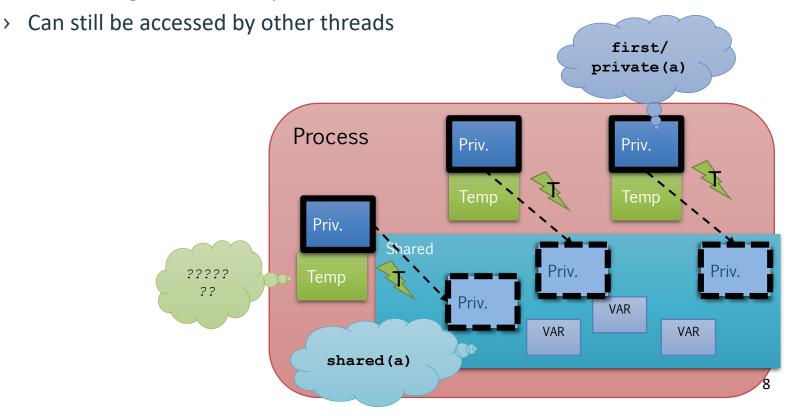
- > Besides synchronization, a barrier has the effect of making threads' <u>temporary view</u> of the shared memory <u>consistent</u>
  - You cannot trust any (potentially modified) shared vars before a barrier
  - Of course, there are no problems with private vars
- > ..what???





# The OpenMP memory model

- > Shared memory with relaxed consistency
  - Threads have access to "a place to store and to retrieve variables, called the memory"
  - Threads can have a <u>temporary view</u> of the memory
    - > Caches, registers, scratchpads...



A bit of architecture...



## Caches in a nutshell

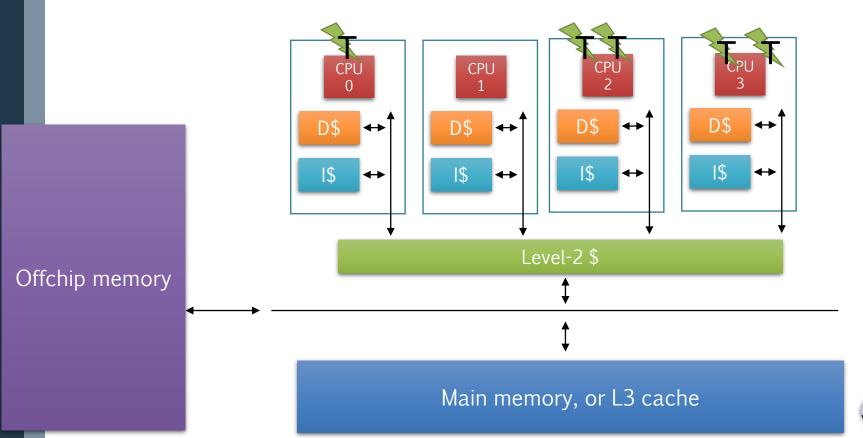
- > A quick memory connected to the core processor
  - ..and to the main memory
  - Few KB of data
- > (If any,) caches are a pure hardware mechanism
  - Used to store a copy mostly accessed data
  - To speedup execution even by 10-20 times
  - Istruction caches/Data caches
- > They perform their work automatically
  - And transparently
  - Poor or no control at all at application level
  - Extremely dangerous in multi- and many-cores



## **Caches**

#### eng.wikipedia.org

A cache is a hardware or software component that stores data so future requests for that data can be served faster; the data stored in a cache might be the result of an earlier computation, or the *duplicate of data stored elsewhere.* 





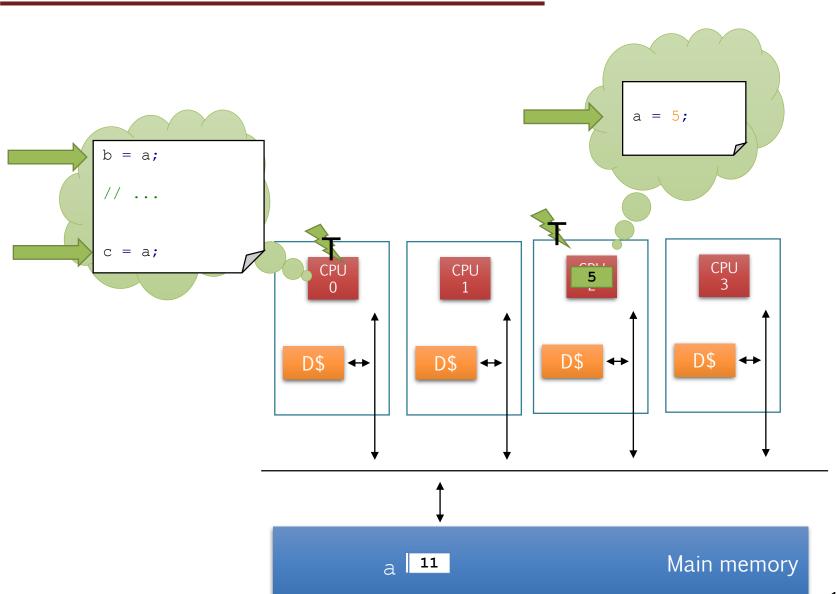
# The catch(es)

- Caches are power hungry
  - Some embedded architectures do not have D\$
- > They are not suitable for critical systems
  - E.g., BOSCH removed D\$s
- > Hardware mechanism, poor control on them
  - Flush command (typically, all cache)
  - Color cache (assign to threads)
  - Prefetch (move data before it's actually needed)

**Coherency** problem in multi/many-cores!!

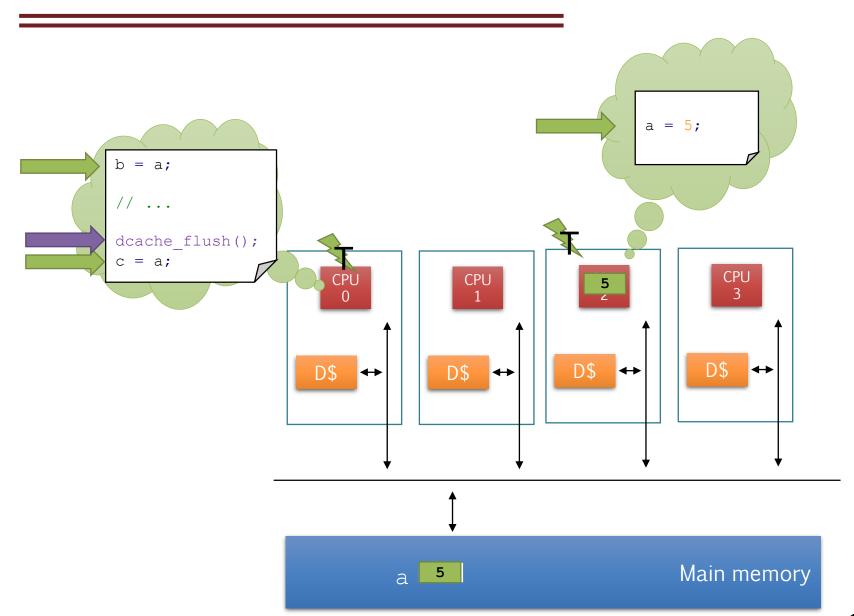


# An example: read stale data



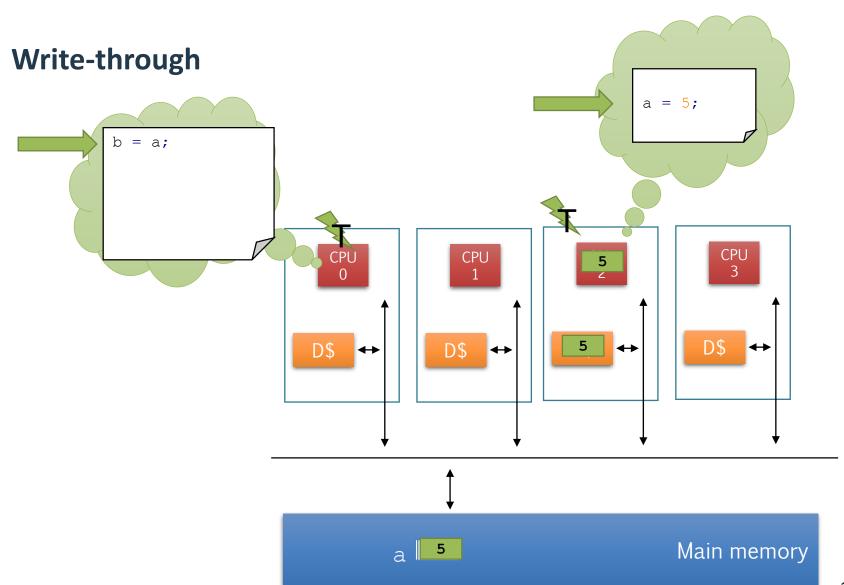


# An example: read stale data



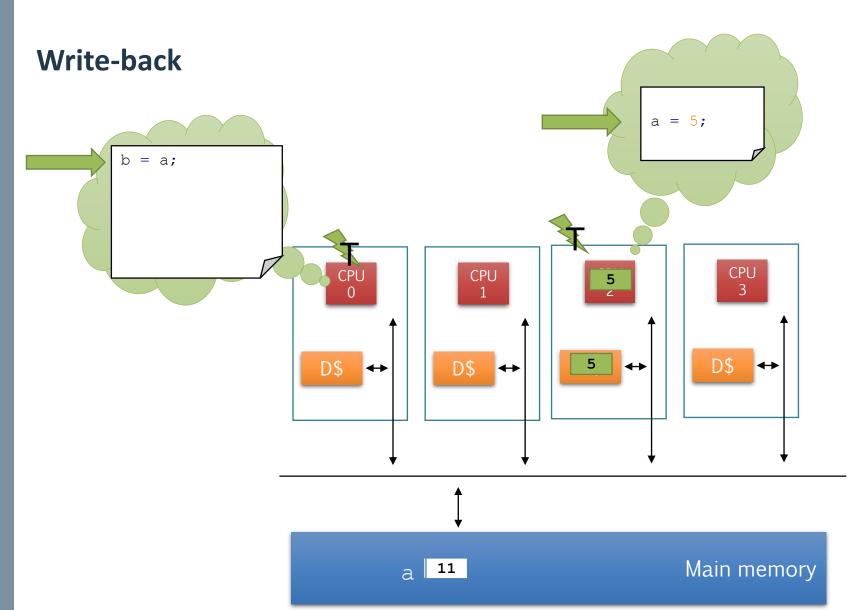


# An(other) example: \$ writing policies



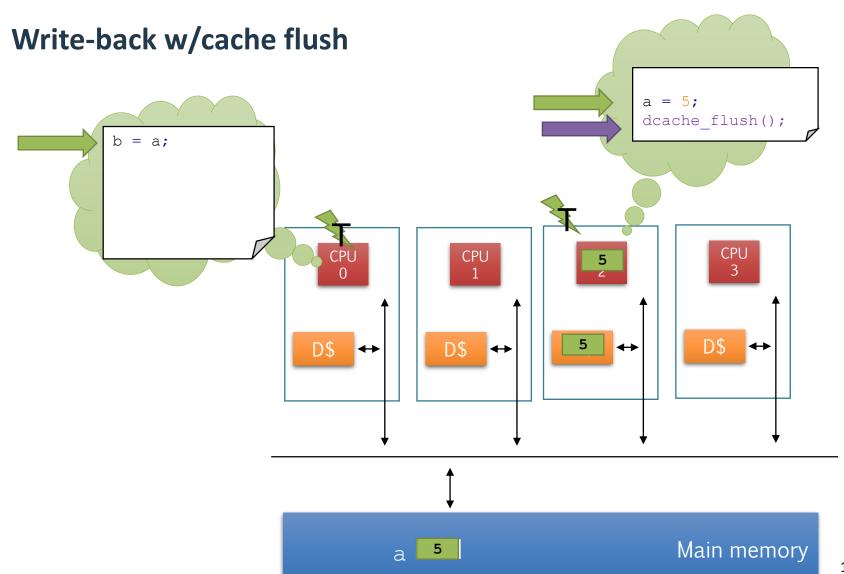


# An(other) example: \$ writing policies





# An(other) example: \$ writing policies





## The flush directive

```
#pragma omp flush [(list)] new-line
```

- > Binding thread set is the encountering thread
  - More "relaxed"
- "It executes the OpenMP flush operation"
  - Makes its <u>temporary view of the shared memory</u> consistent with other threads
  - "Calls to dcache\_flush()"
- > Enforces an order on the memory operations on the variables specified in list



### Semantics: barrier vs flush

#### #pragma omp barrier

- > Joins the threads of a team
- > Applies to all threads of a team
- > Forces consistency of threads' temporary view of the shared memory

#### #pragma omp flush

- > Applies to one thread
- > Forces consistency of its temporary view of the shared memory
- > Much lighter!



## **OpenMP software stack**



- > Multi-layer stack
  - Engineered for portability

```
User code

a = 5;

#pragma omp flush
```

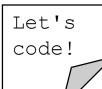
```
OpenMP runtime void GOMP_flush() {
    dcache_flush();
}
```

```
Operating System void dcache_flush()
{
    asm("mov r15, #1");
}
```

Hardware D\$



## How to run the examples



> Download the Code/ folder from the course website

- Compile
- > \$ gcc -fopenmp code.c -o code

- > Run (Unix/Linux)
- \$ ./code
- > Run (Win/Cygwin)
- \$ ./code.exe



#### References



- > "Calcolo parallelo" website
  - http://hipert.unimore.it/people/paolob/pub/Calcolo Parallelo/
- > My contacts
  - paolo.burgio@unimore.it
  - http://hipert.mat.unimore.it/people/paolob/
- > Useful links
  - http://www.google.com
  - http://www.openmp.org
  - <a href="https://gcc.gnu.org/">https://gcc.gnu.org/</a>
- > A "small blog"
  - http://www.google.com